
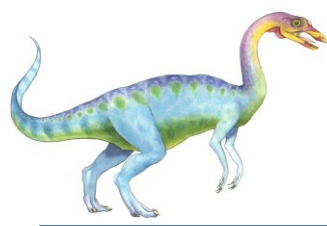


# Memory Management





# Basic Concepts

---

## Physical memory

- Hardware memory shared by all programs
- Organized for performance (memory hierarchy, cache)
- Vary in size on each machine

## Logical memory

- Abstraction of contiguous memory address space
- Simple memory model for programming
- Same size for every process





# OS tasks for Memory Management


T

- Provide an abstraction of logical memory for programs
  - Chapter 8
- Provide a virtual memory that is larger than physical memory.
  - Chapter 9
- Others
  - Protection from other processes
  - Memory sharing between processes
  - Memory-mapped I/O



# Chapter 8: Memory- Management Strategies


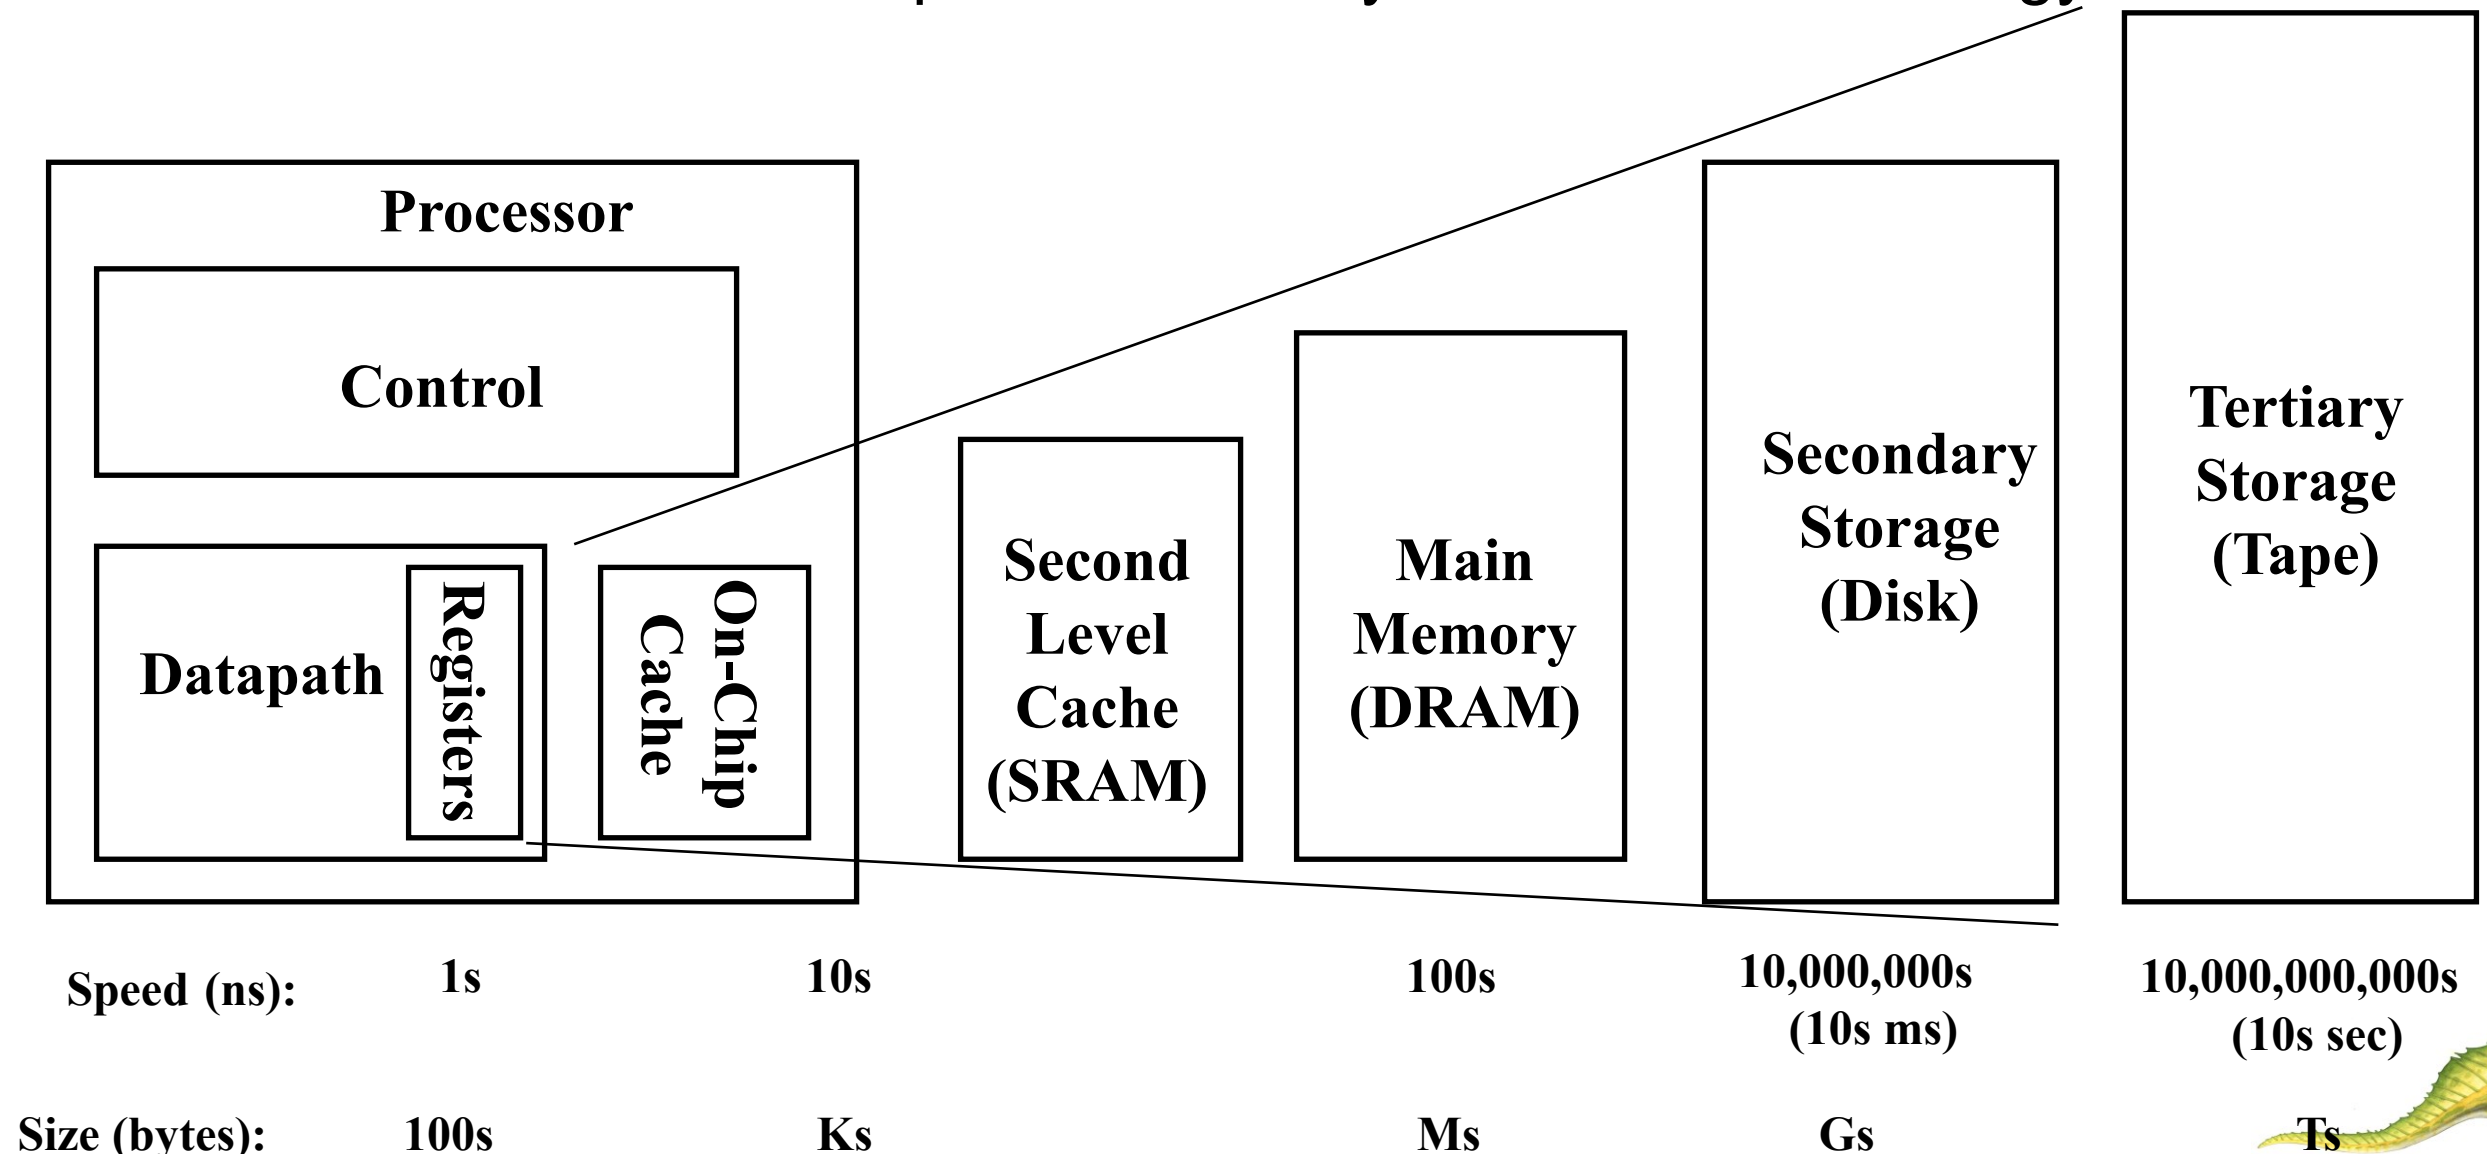
---





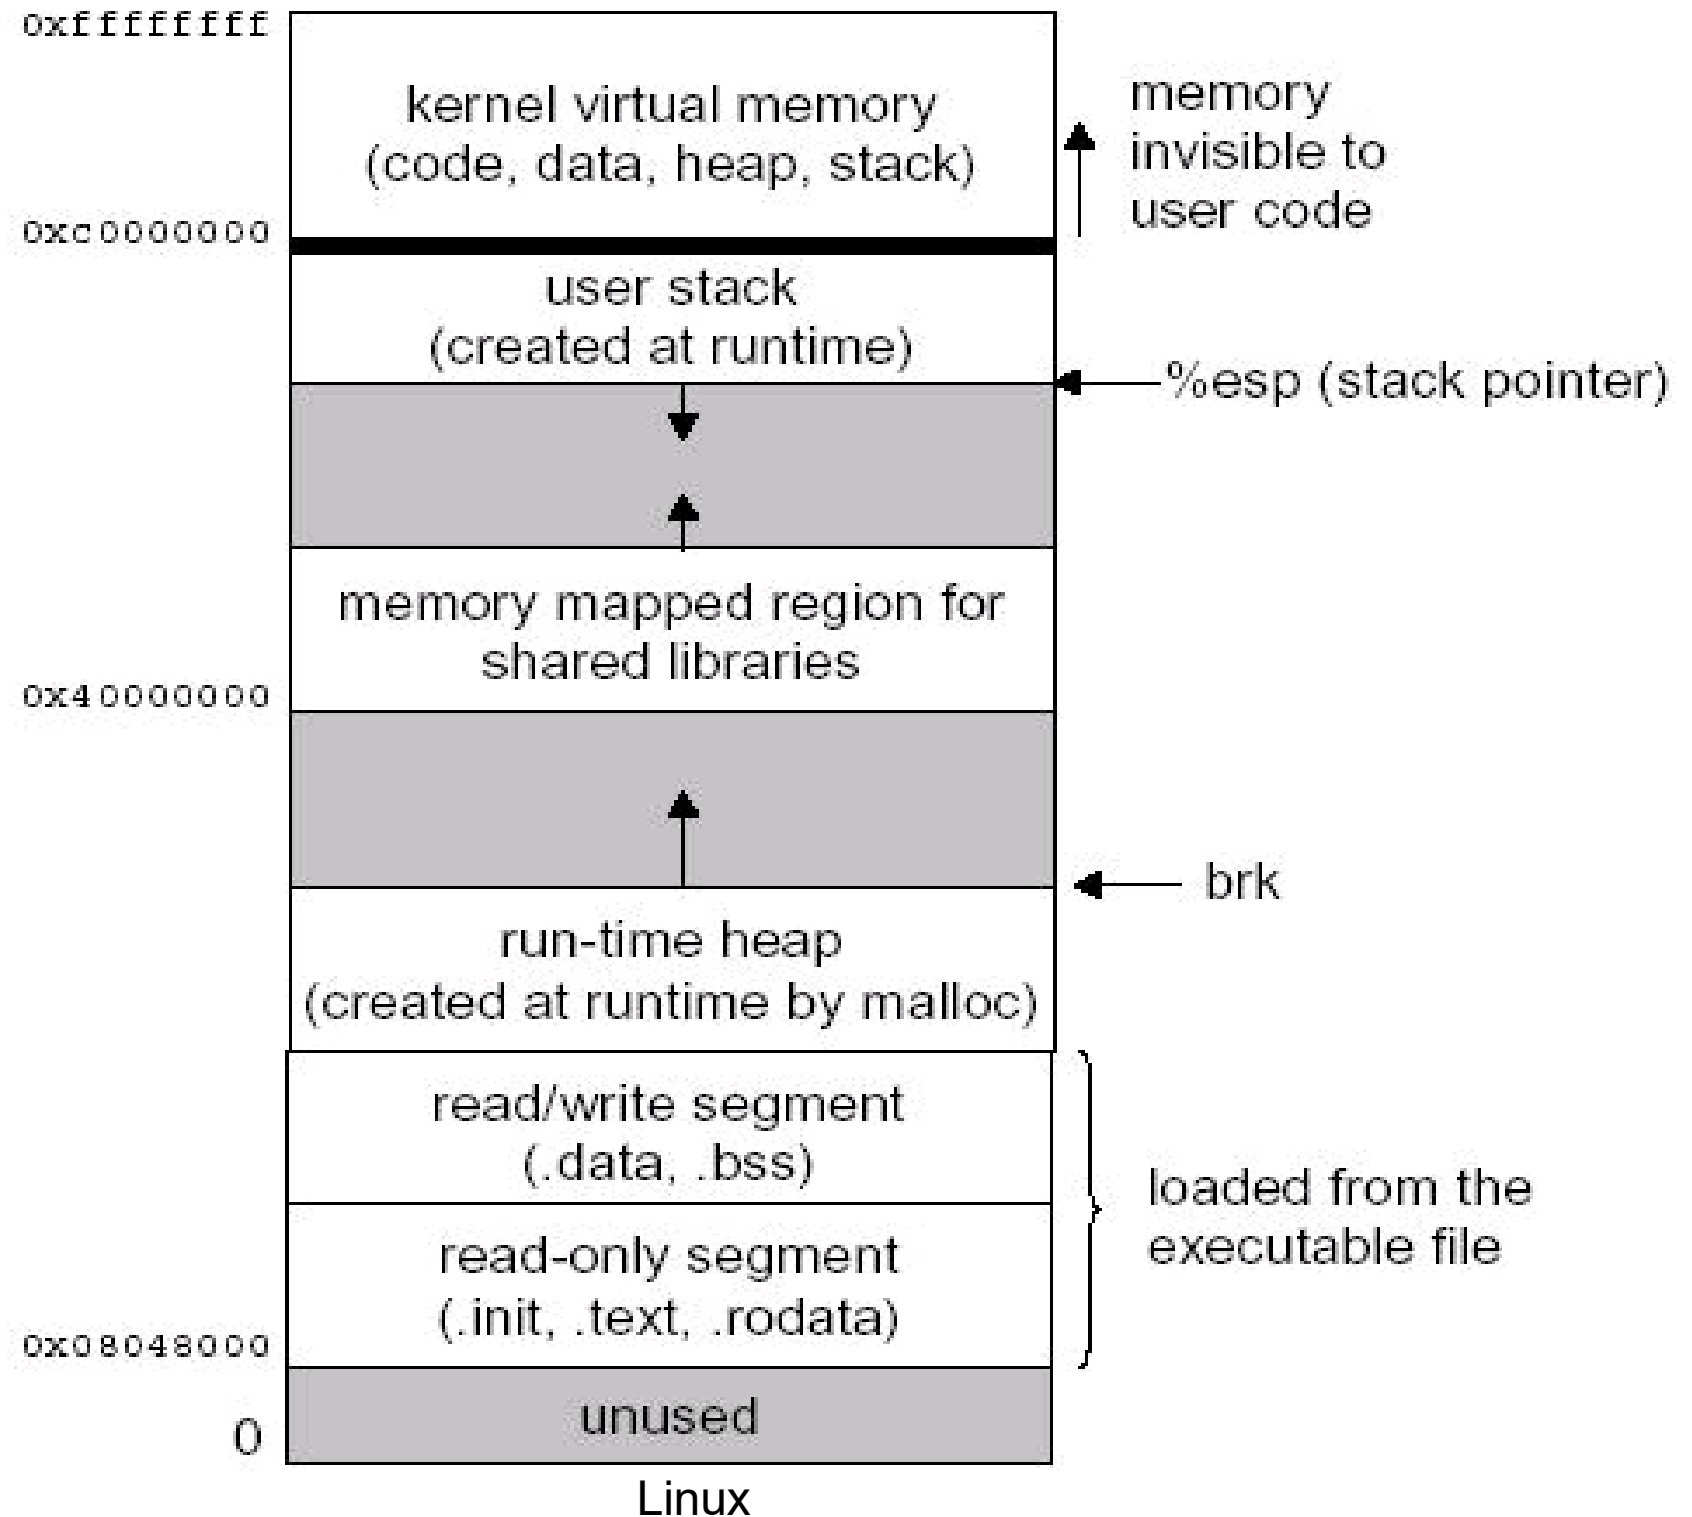

# Memory Hierarchy of a Modern Computer System

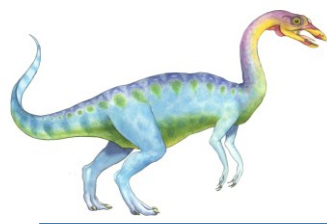
- By taking advantage of the principle of **locality**:
  - Present the user with as much memory as is available in the cheapest technology.
  - Provide access at the speed offered by the fastest technology.





# Process Memory Layout





# Sample C Program

---

```
#include <math.h>

int global;

int f(int farg)
{
    int flocal;

}

main()
{
    int local;
    int *dynamic;
    dynamic = malloc(10, sizeof(int));
    f(1);
    local = sqrt(2.0);
}
```





# Addresses

Size of executable file = 314740 bytes

address of f()	1020c	}	Code segment
address of main()	10244		
address of sqrt()	10358		
address of printf()	13128		
address of malloc()	31218	}	Data segment
address of global	5004c		
value of dynamic	50068	}	Heap segment
		↑↓	
address of f()'s flocal	ffbefb4c	}	Stack segment
address of f()'s farg	ffbefba4		
address of main()'s local	ffbefbc4		





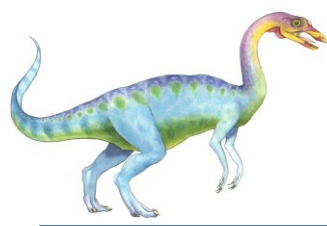


# Logical vs. Physical Address Space


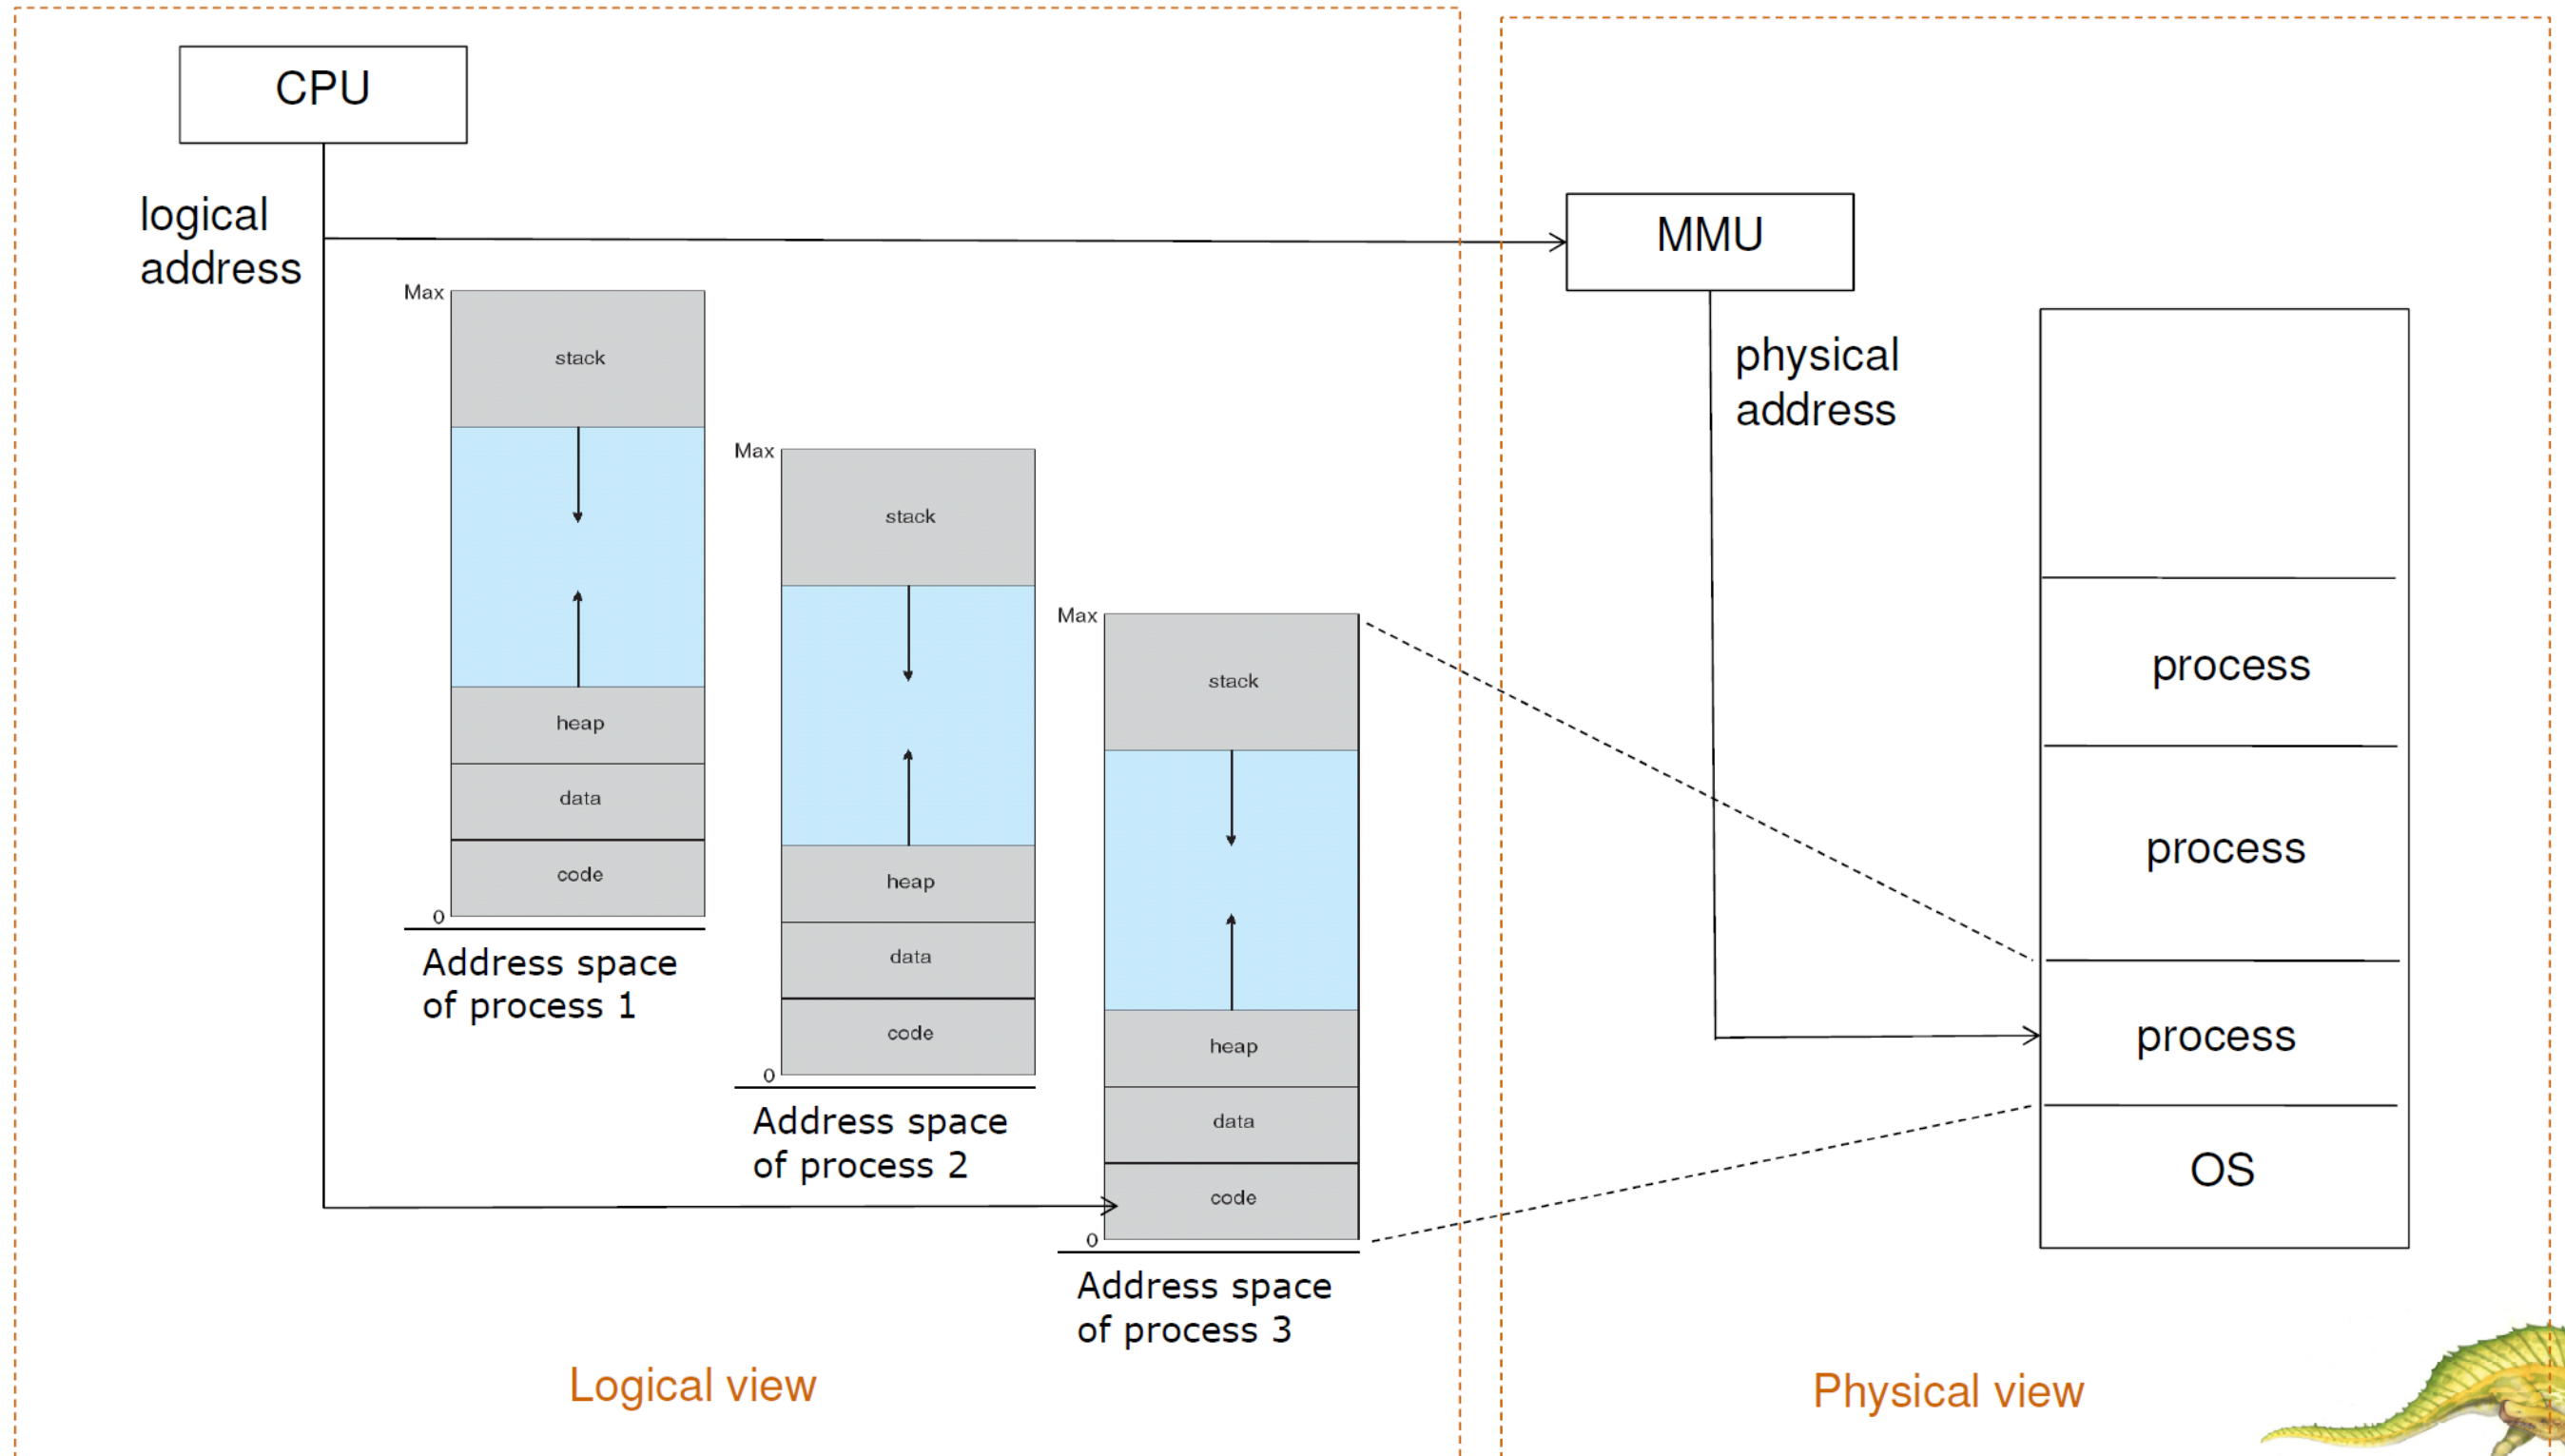
---

- **Logical address** – address in program instructions that CPU executes; also referred to as **virtual address**
- **Physical address** – address used by memory unit to access physical memory (RAM)
- Every logical address must be translated into physical address by Memory Management Unit (MMU).
- The user program deals with logical addresses; it never sees the real physical addresses






# Address Translation









# MEMORY ALLOCATION







## Contiguous Allocation




Cannot find  
space of new  
process






## Contiguous Allocation




Allocate one large piece of memory for each process.

Cannot find space of new process




## Paging



Allocate many small pieces of memory for each process.

Can find space of new process





# Paging

- Physical memory is divided into fixed-sized blocks called **frames**
- Logical memory is divided into blocks of same size called **pages**.
- Size of frames and pages is power of 2, e.g. 512, 1K, 2K, **4K**, 8K bytes
- OS keeps track of all **free frames**
- To run a program of size ***n*** pages, need to find ***n*** free frames and load program
- Use a **page table** is used to translate logical to physical addresses





# Address Translation Scheme

Logical Address is divided into:

**Page number ( $p$ )** – used as an index into a *page table* which contains base address of each page in physical memory


**Page offset or displacement ( $d$ )** – combined with base address to define the physical memory address that is sent to the memory unit



For given logical address space  $2^m$  and page size  $2^n$

		page	offset
0	a	0	0
1	b		1
2	c		2
3	d		3
4	e	1	0
5	f		1
6	g		2
7	h		3
8	i	2	0
9	j		1
10	k		2
11	l		3
12	m	3	0
13	n		1
14	o		2
15	p		3

logical memory





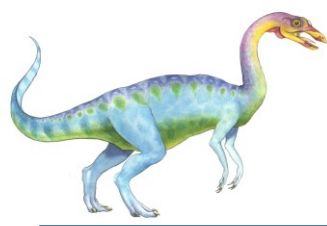
# Example: page no. and offset

Address width	32 bits
Page size	4K bytes
Offset	12 bits
Page number	20 bits

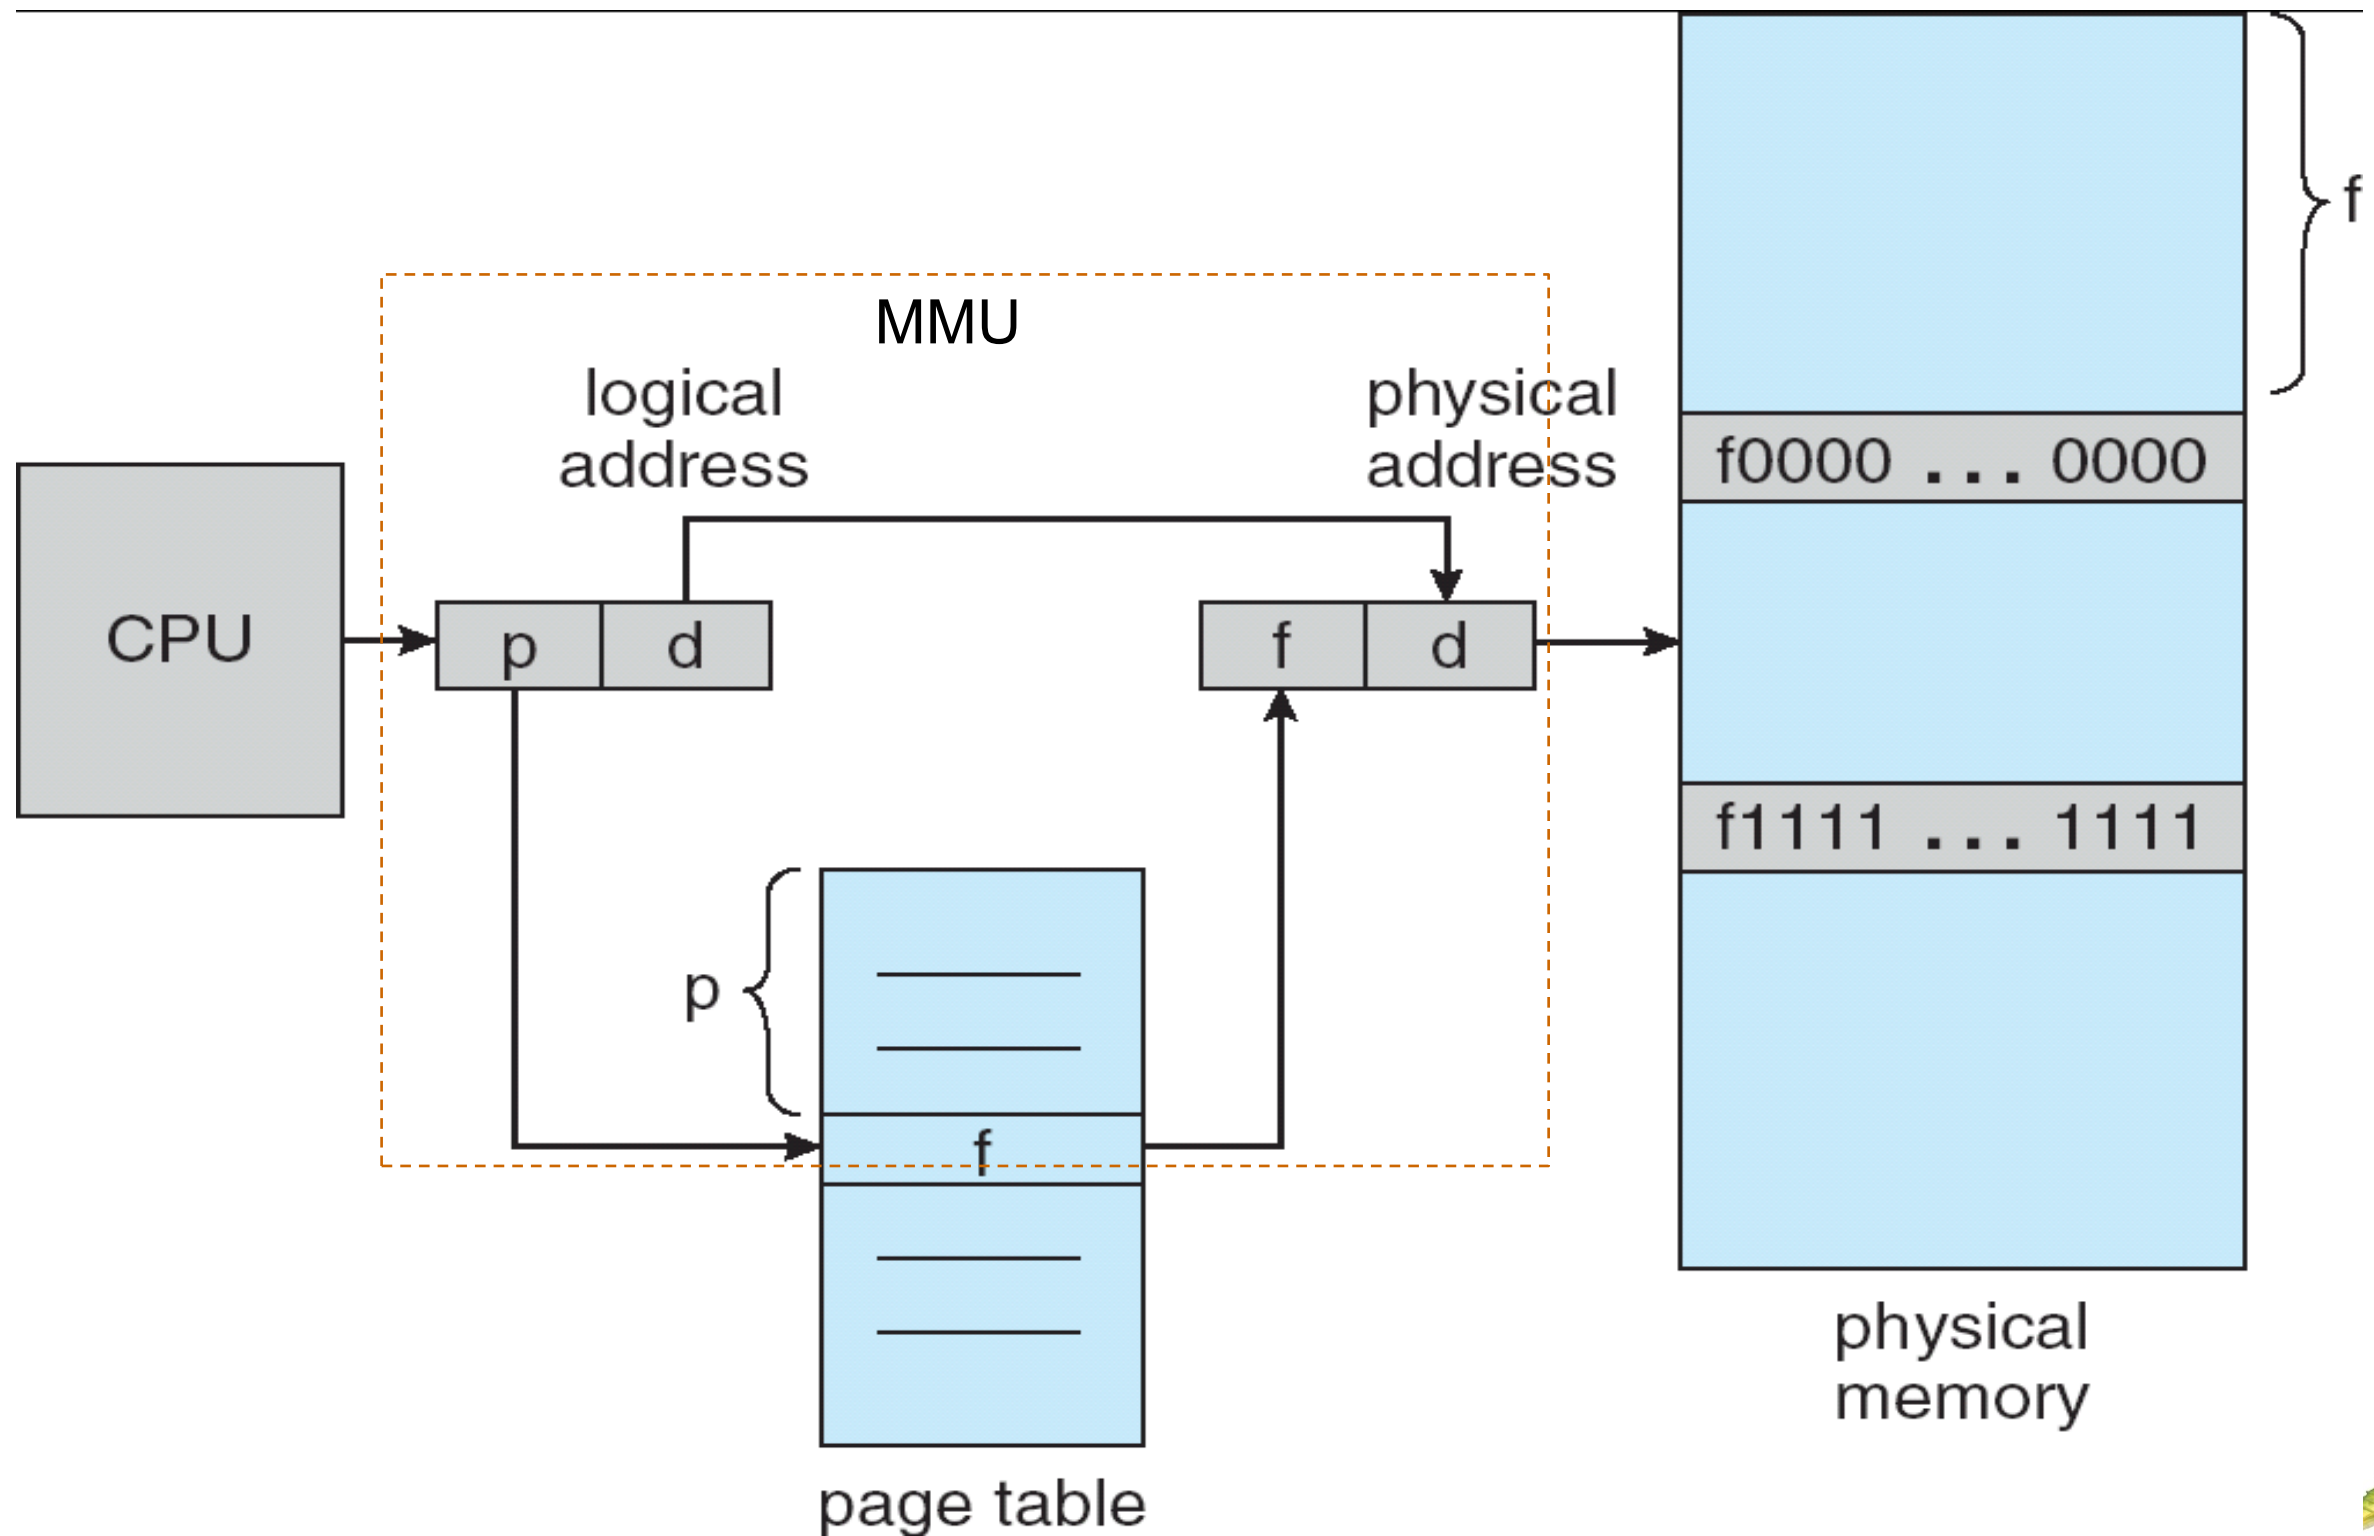
	page	offset
f()	10	20c
main()	10	244
sqrt()	10	358
printf()	13	128
malloc()	31	218
global	50	04c
flocal	ffbef	b4c
farg	ffbef	ba4








# Paging Hardware





# Paging Model of Logical and Physical Memory




logical  
memory

0	1
1	4
2	3
3	7

page table

frame  
number

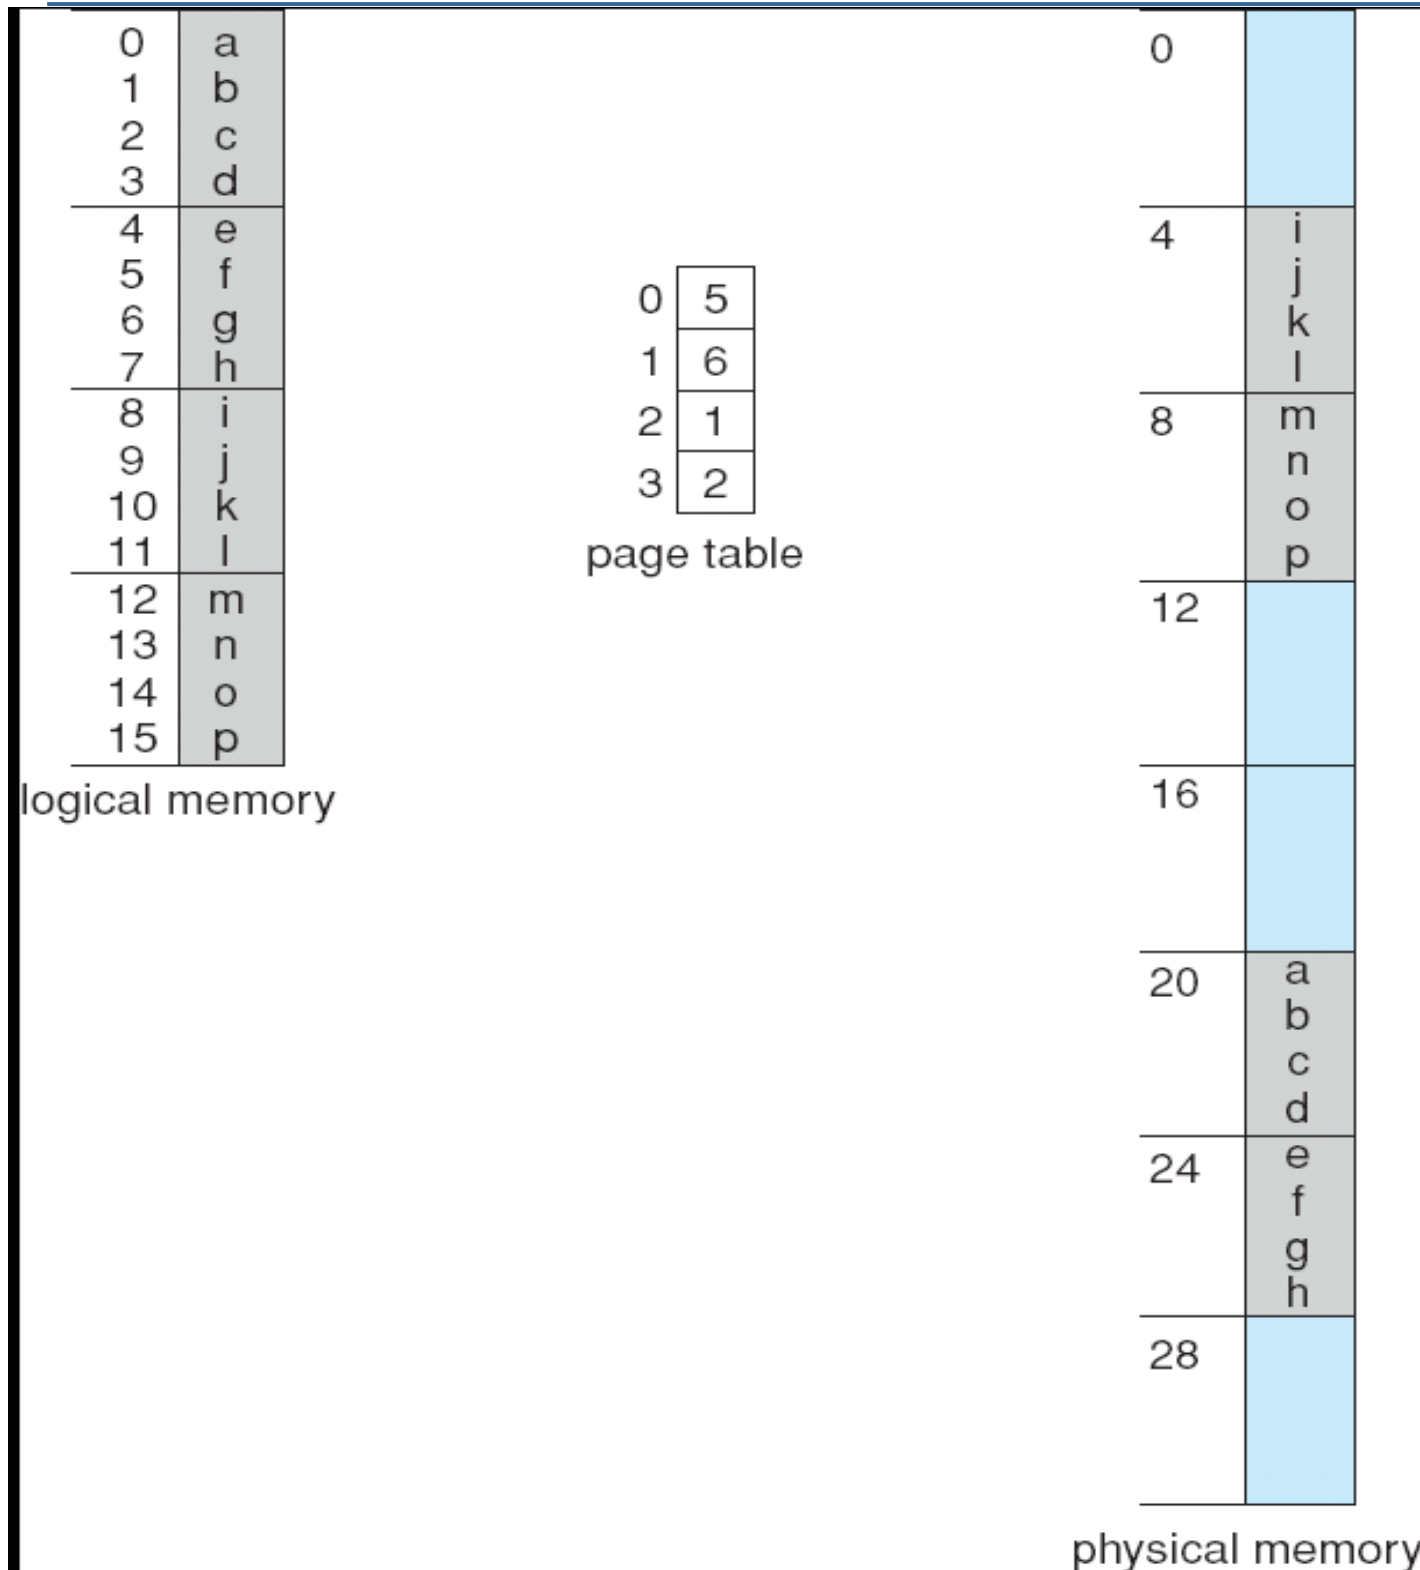


physical  
memory





# Paging Example



16-byte logical memory

4-byte page size → 4 pages

4-bit address → 2-bit page number, 2-bit offset

32-byte physical memory

4-byte frame size → 8 frames

5-bit address → 3-bit frame number, 2-bit offset

Ex.

Logical address = 1001

Page number = 10

Frame = 001

Offset = 01

Physical address = 00101

Data = 'j'






# Translation Look-Aside Buffer

- Page table is kept in main memory
- **Two memory access problem:** Every data/instruction access requires two memory accesses. One for the page table and one for the data/instruction.
- The problem can be solved by using a special hardware cache called **translation look-aside buffers (TLBs)**
- Keep recently used page table entries
- Address translation (p, d)
  - Search for p in TLB
  - If p is found (TLB hit), get frame no. out
  - Otherwise (TLB miss), get frame # from page table in memory

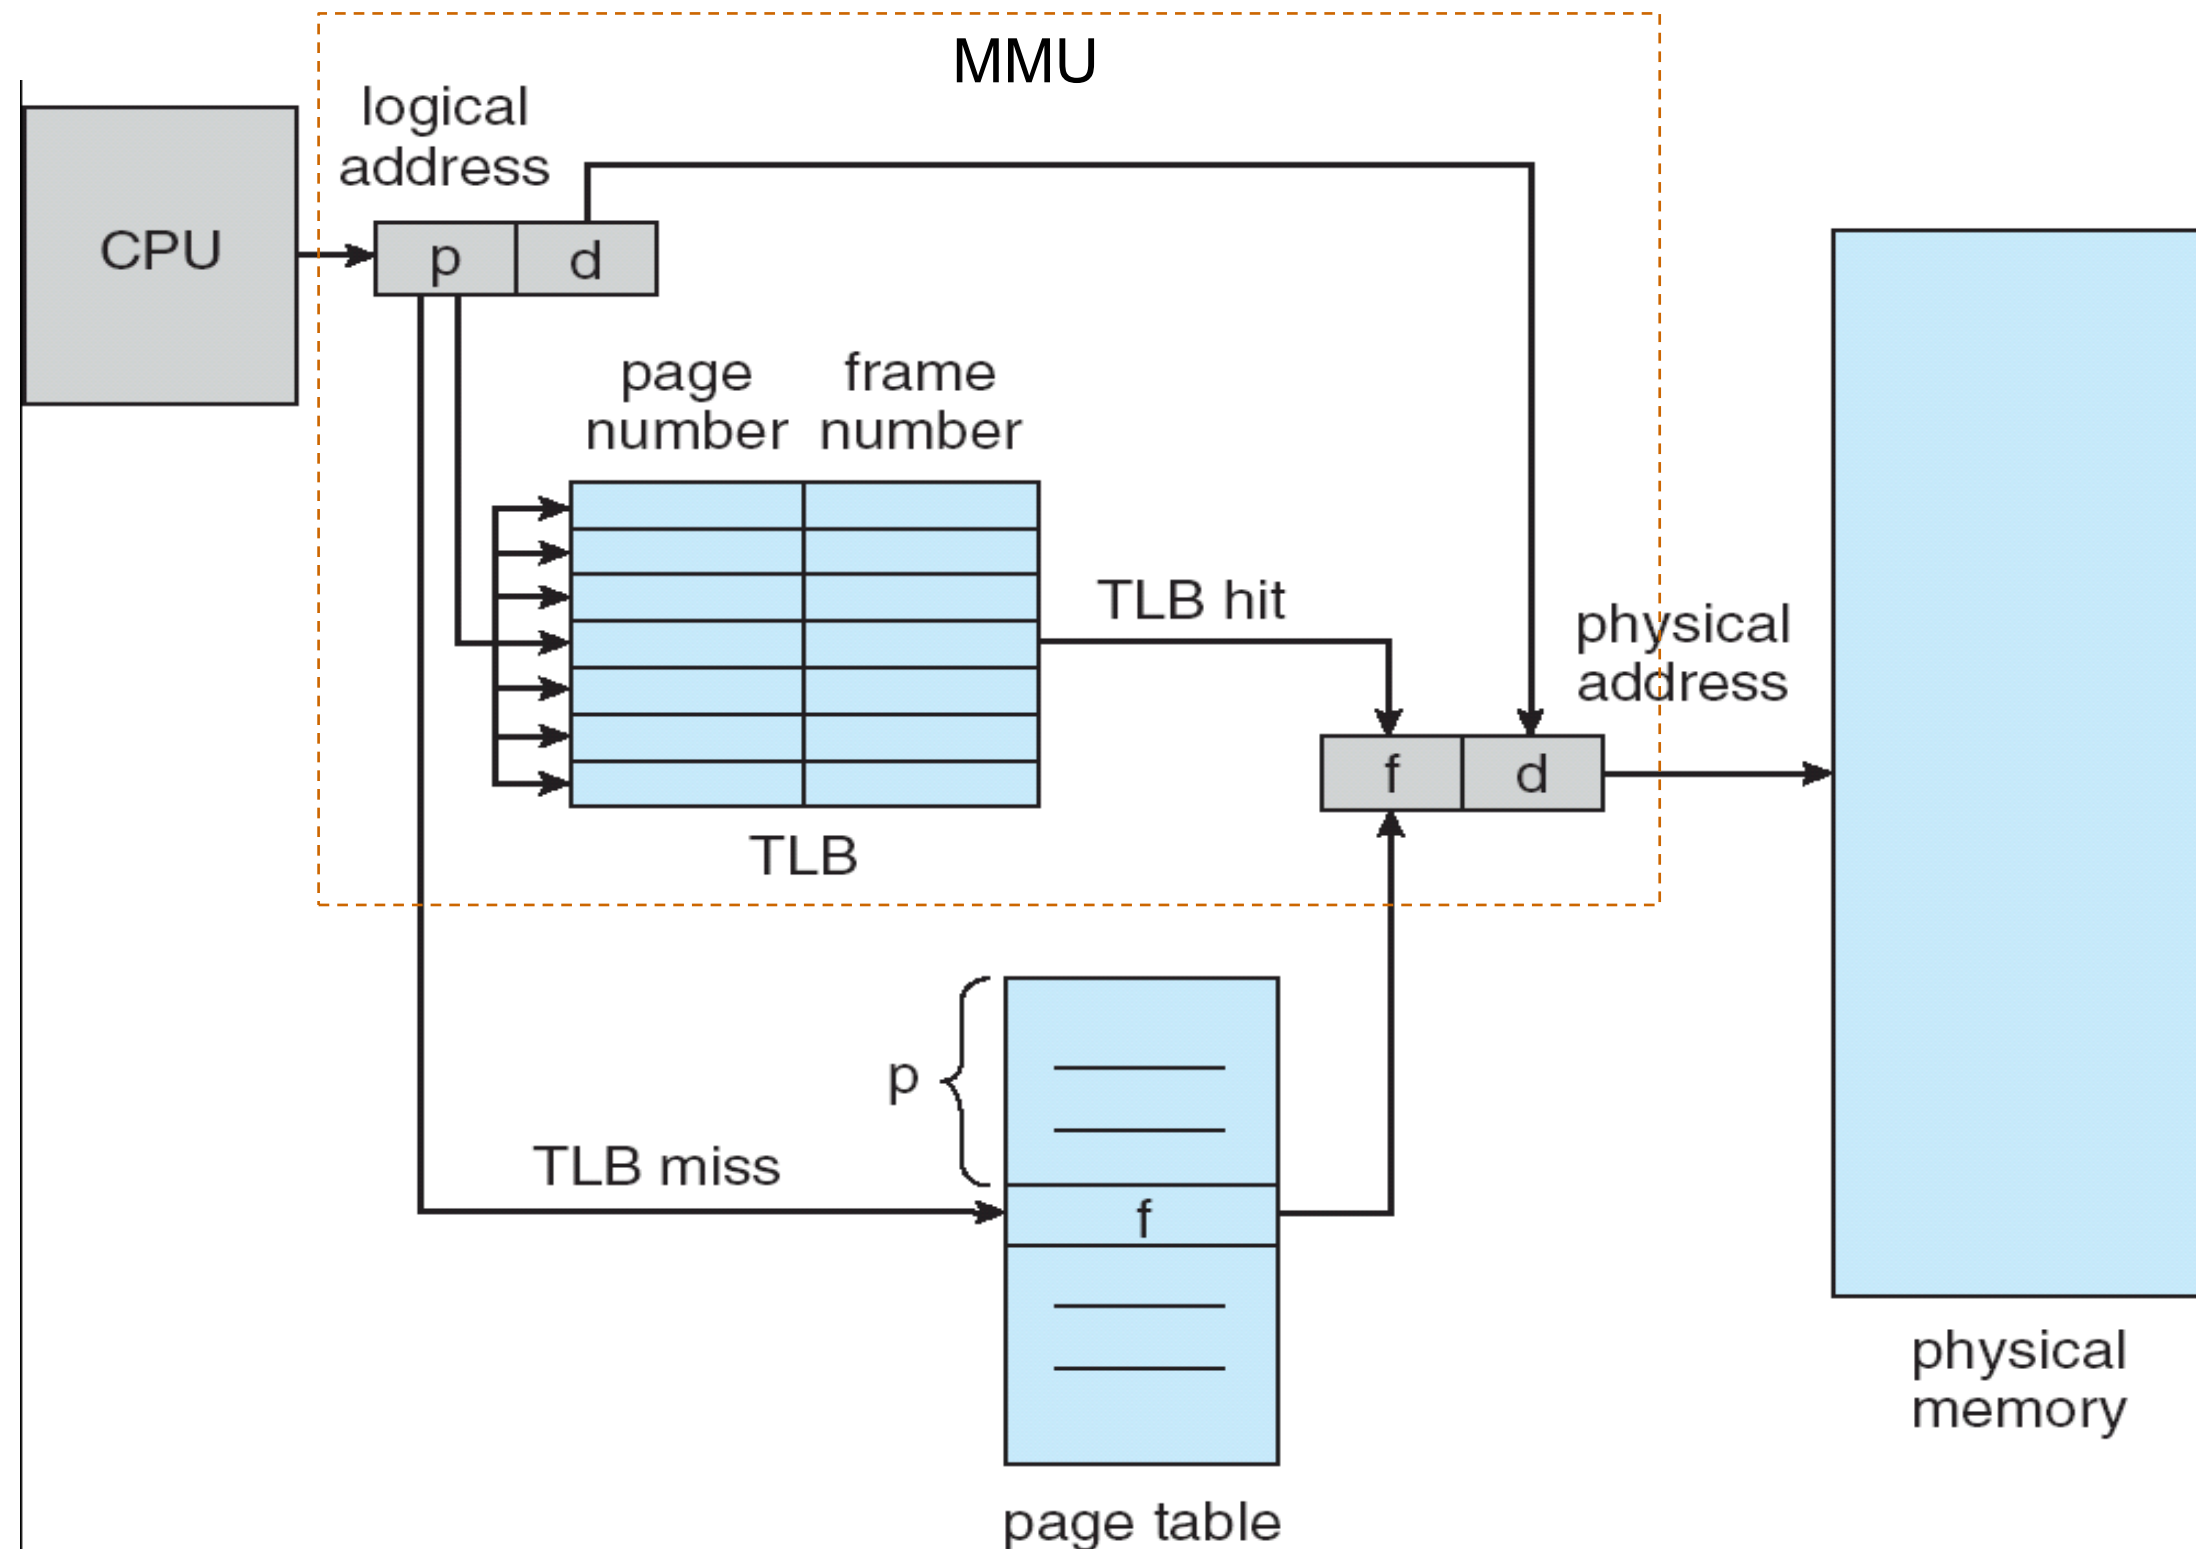
page number	frame number

TLB





# Paging Hardware With TLB





# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
- Memory access time =  $m$  time unit
- Hit ratio – percentage of times that a page number is found in TLB
- Hit ratio =  $\alpha$
- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (\varepsilon + m) \alpha + (\varepsilon + 2m)(1 - \alpha) \\ &= \varepsilon \alpha + m \alpha + \varepsilon + 2m - \varepsilon \alpha - 2m \alpha \\ &= 2m - \alpha m + \varepsilon \end{aligned}$$







# Structure of the Page Table

- For 32-bit virtual address space with 4 KB page size, page table may have up to 1 million ( $2^{32}/2^{12}$ ) entries (per process!)
- Intel Itanium 2 processor supports 64-bit virtual address space and 50-bit physical address space with different page sizes (4kB, 8kB, 16kB, 256kB, 1MB, 4MB, 16MB, 64MB and 256MB)
- Current operating systems support up to 50-bit virtual address space
- A lot of pages between heap and stack is unused, but page table entries are still allocated for them (if page table is an array)
  
- Aim: reduce size of page table
- Common page table structures
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables





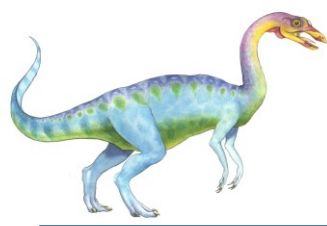
# Hierarchical Page Tables

---

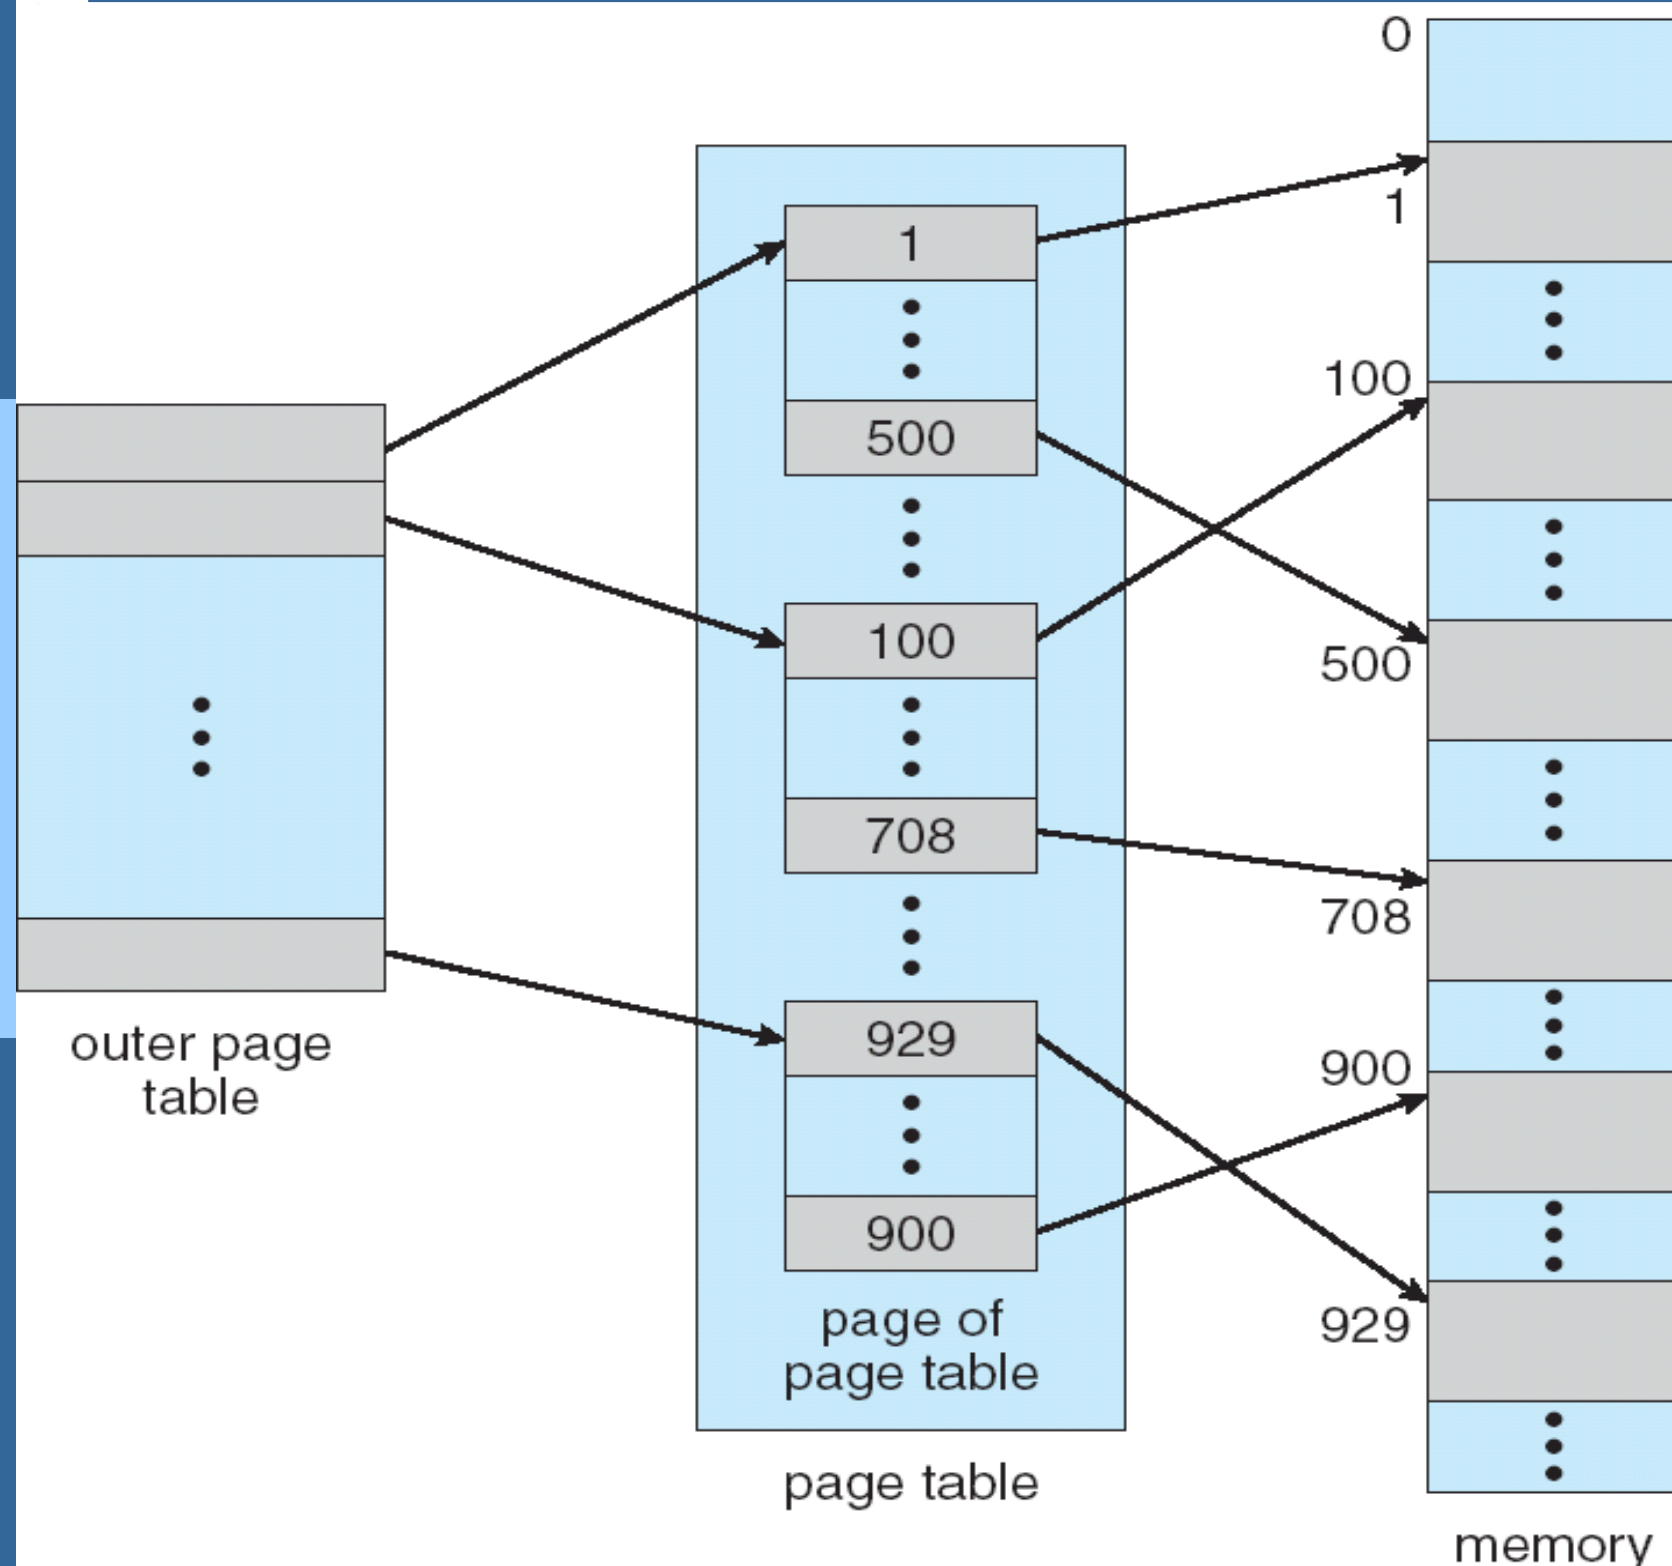
- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- Page table is also paged







# Two-Level Page-Table Scheme




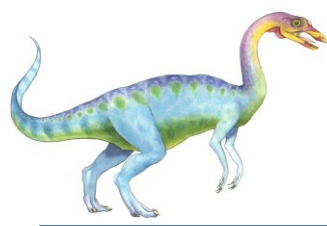
## Outer page table

- contains all entries for the whole logical address space
- each entry associates to a lot of pages
- reduce number of entries
- entries for unused memory area are null

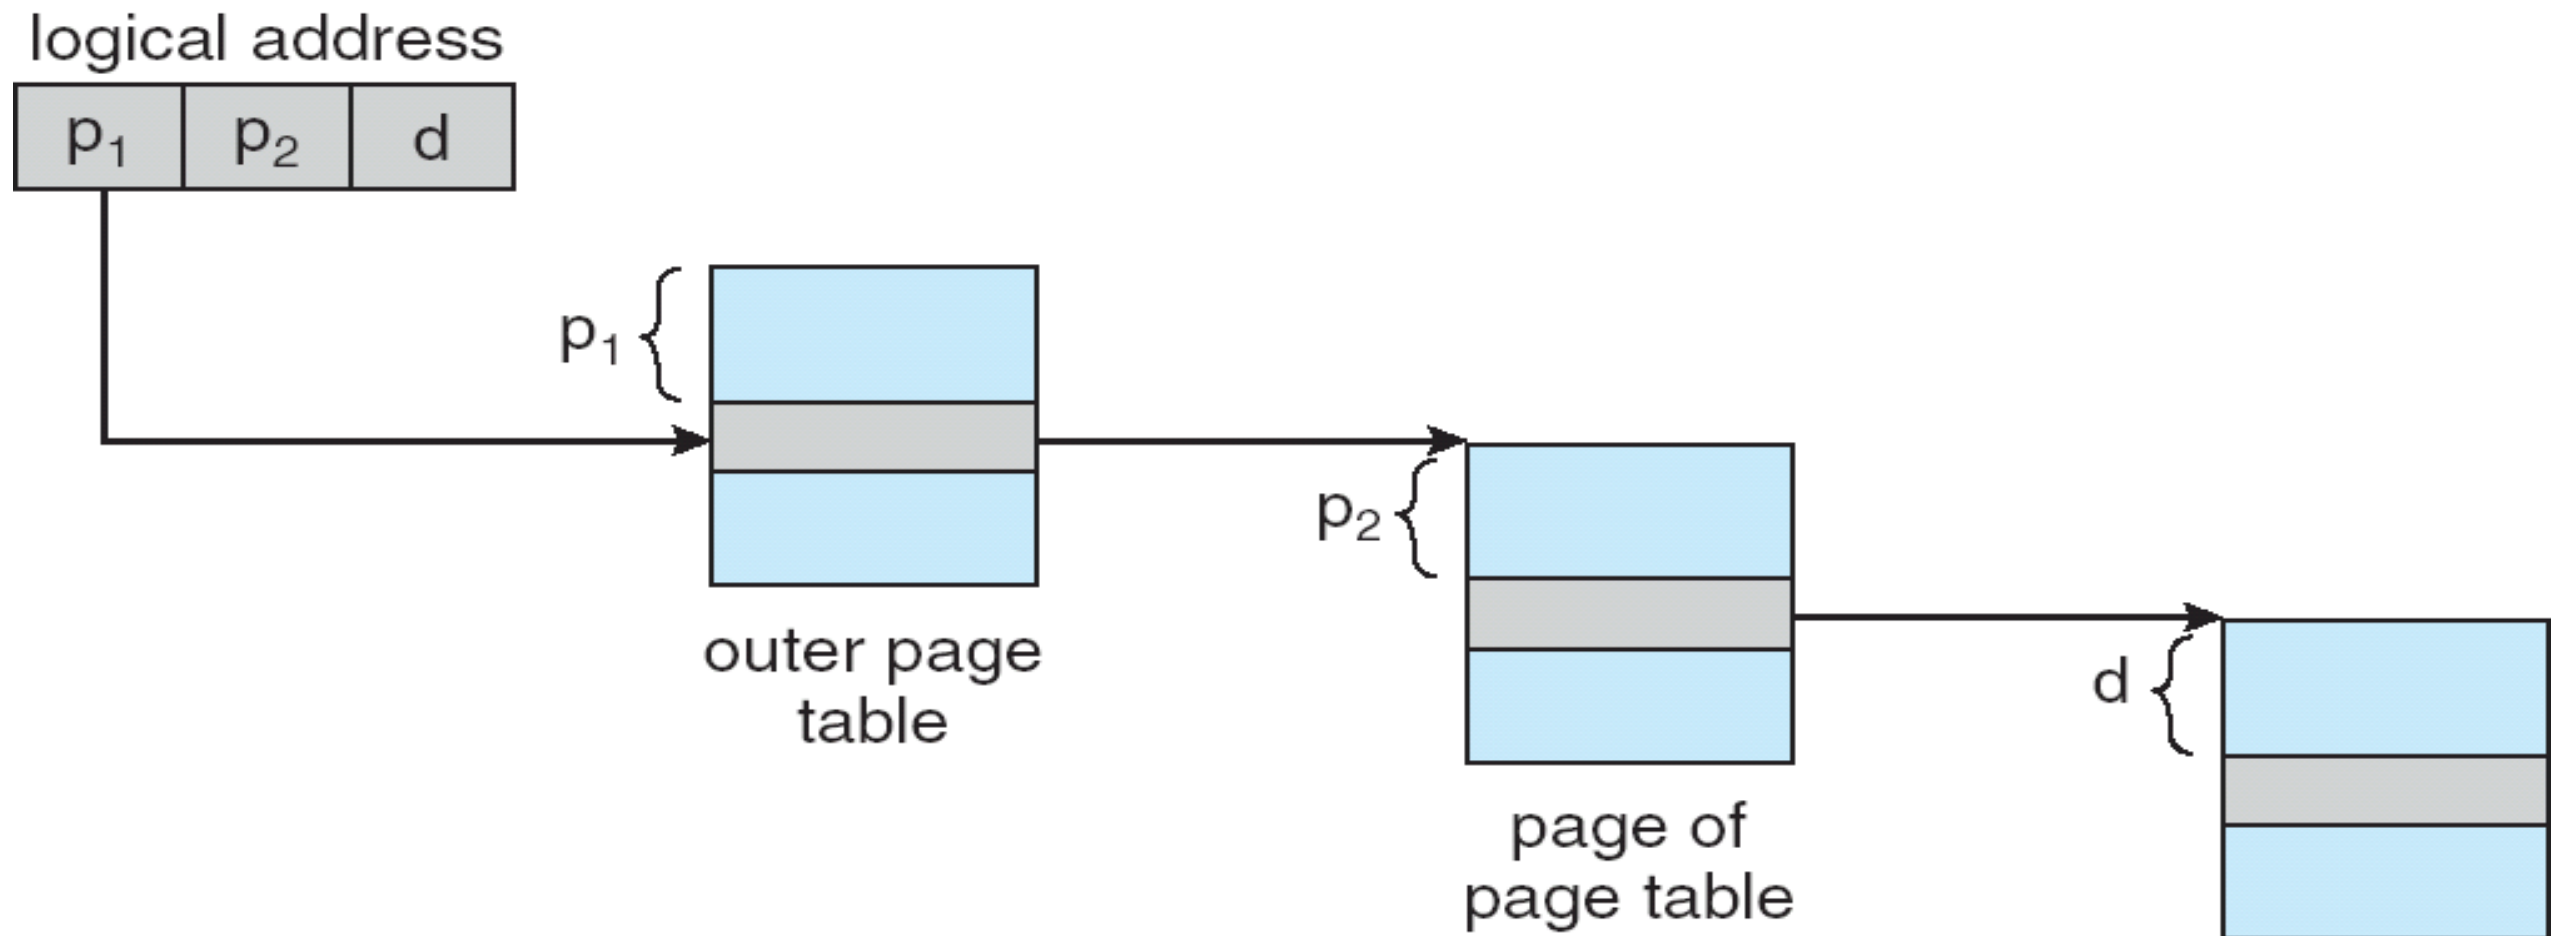
## page table

- contains pages of page table for only used pages
  - dynamically allocated
  - reduce number of entries





# Two-Level Address-Translation





# Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
  - a page number consisting of 20 bits
  - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
  - a 10-bit page number ( $p_1$ )
  - a 10-bit page offset ( $p_2$ )
- Thus, a logical address is:

page number		page offset
$p_1$	$p_2$	$d$
10	10	12

where  $p_i$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table

- Outer page table has 1K entries
- page table contains pages of page table, each has 1K entries





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12





# Three-level Paging in Linux

