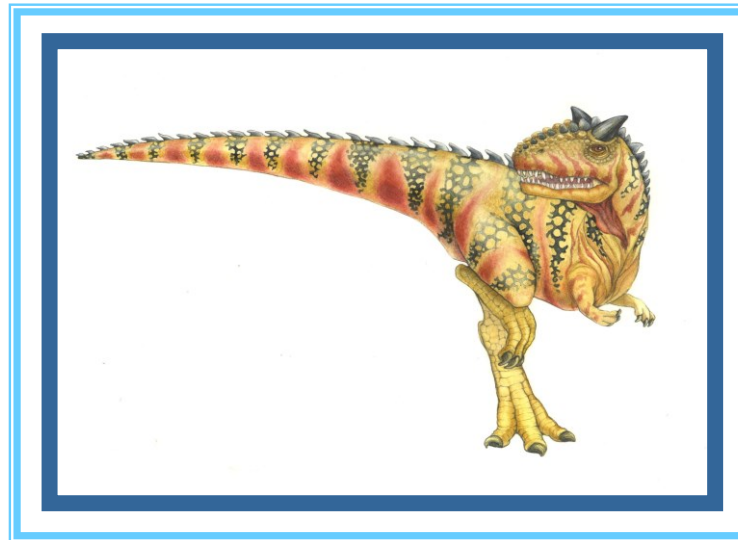
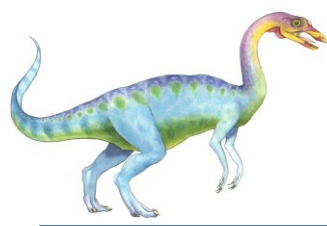


Chapter 9: Virtual-Memory Management





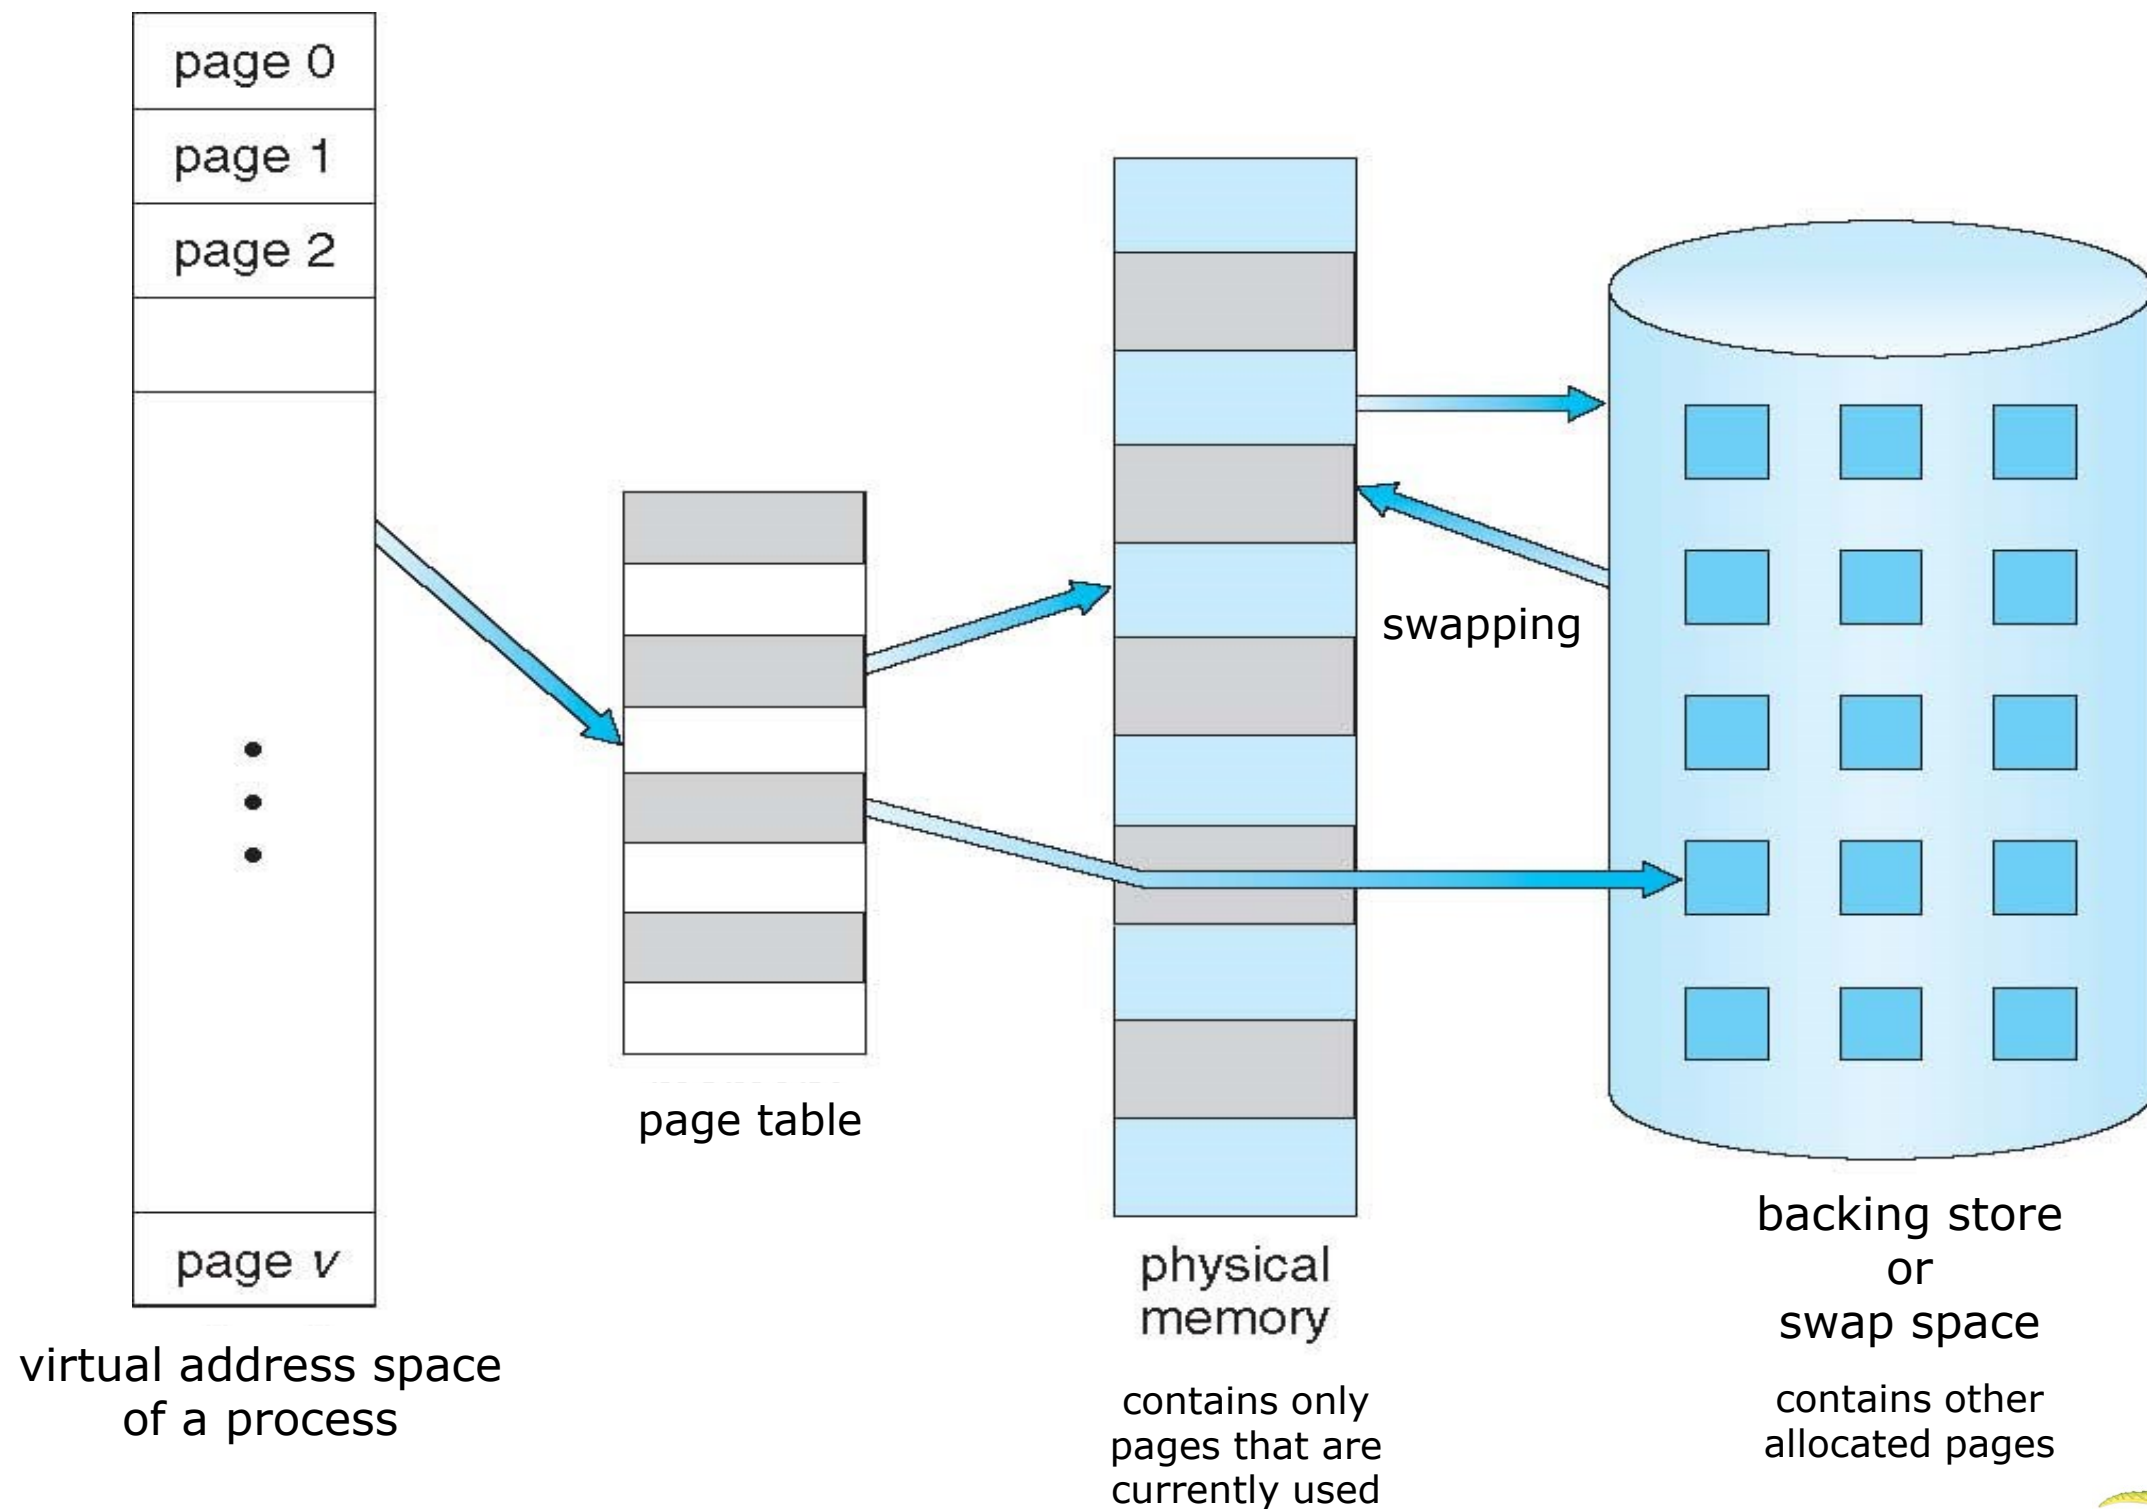
Virtual Memory

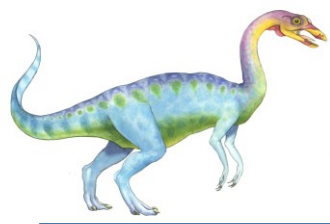
- Goal
 - Provide memory space for all processes that can be much larger than physical memory.
- How
 - Only some part of program code and data need to be in physical memory at a time.
 - ▶ principle of locality
 - The rest of program code and data can be stored in secondary storage (hard disk)
 - ▶ “swap partition” or “swap file”
 - ▶ extend memory hierarchy
 - Implement “Demand Paging” mechanism in OS
 - ▶ similar to “cache miss”
- Other benefits
 - Address spaces can be shared by several processes
 - ▶ shared program code, shared libraries, shared memory (IPC)
 - fast process creation (allocate memory only when needed)
 - memory-mapped file, memory-mapped I/O





Virtual Memory





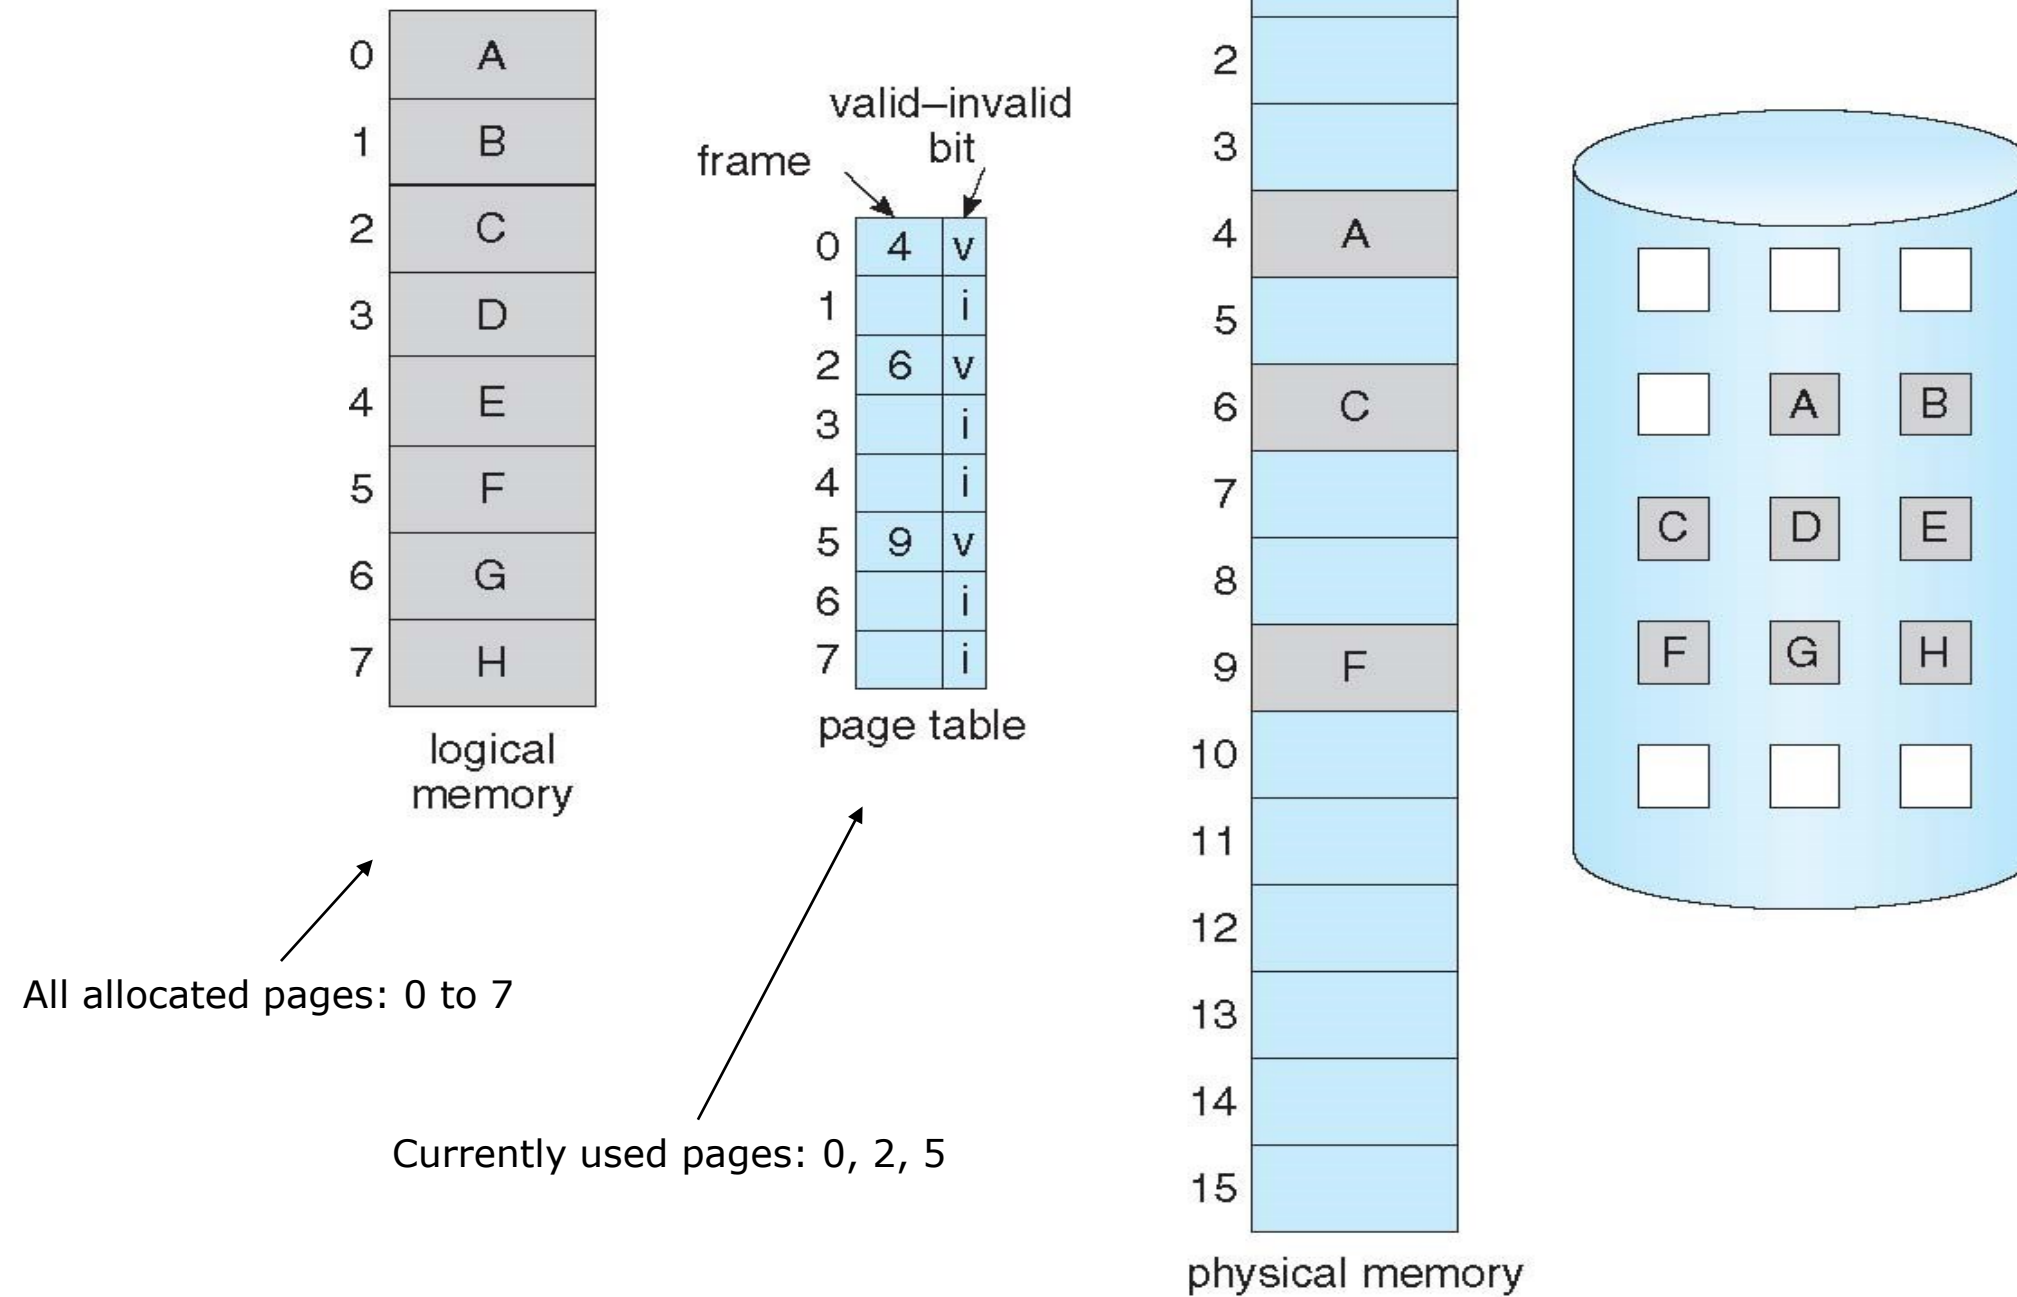
Demand Paging

- Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More processes
- CPU instruction references to anywhere in a page \Rightarrow the page needs to be in physical memory
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
 - no-free-frame \Rightarrow swap





Page table that supports demand paging





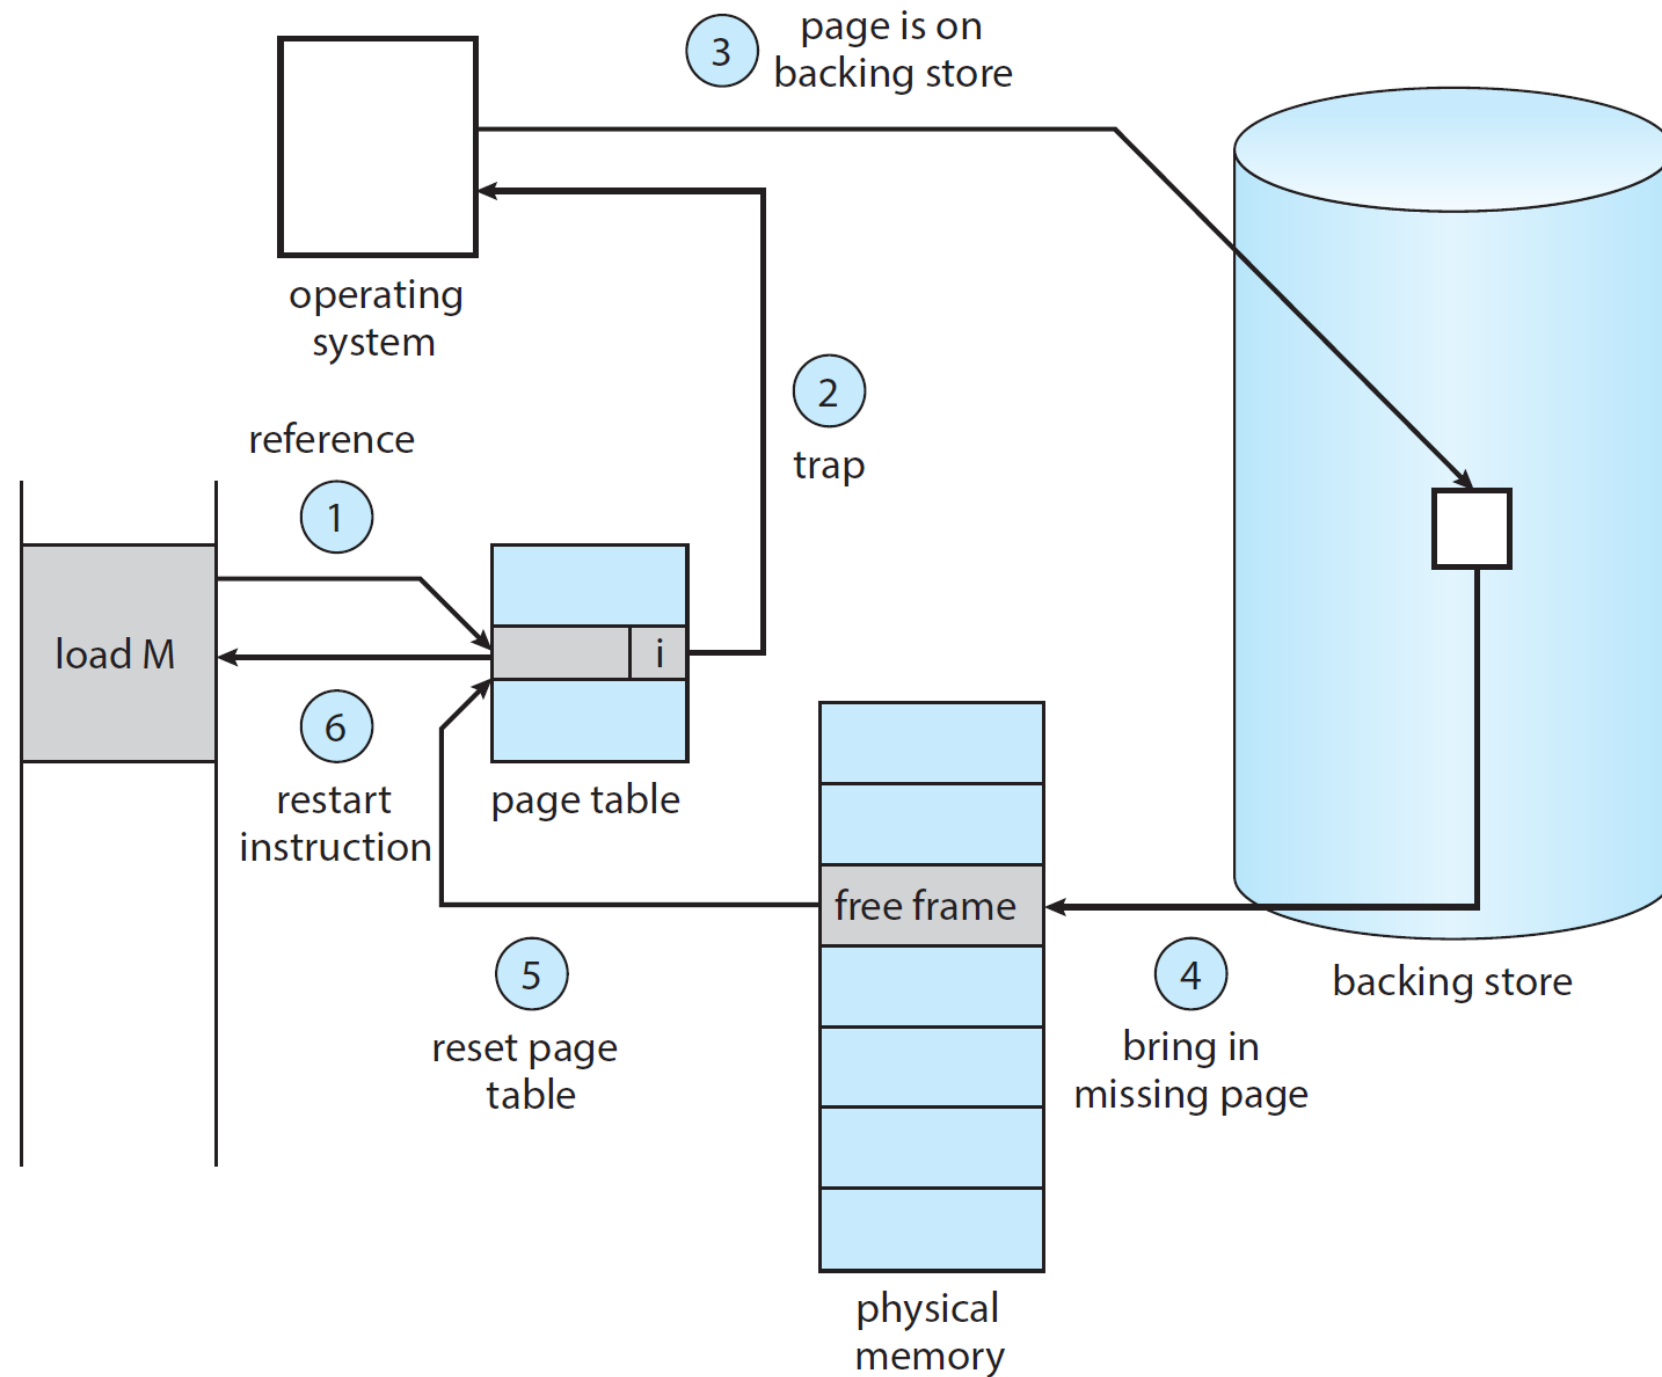
Page Fault

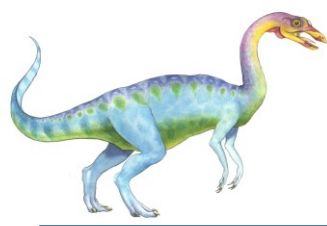
- Page fault is a mechanism to implement Demand Paging
 - Similar to cache miss
- A reference to an invalid page (e.g. first reference to that page) will cause an interrupt to operating system:
 - page fault interrupt (generated by MMU)
 - page fault handler (code in OS kernel)
- 1. Operating system looks at another table to decide:
 - Invalid reference → abort
 - Just not in memory → continue
- 2. Get empty frame
- 3. Swap page into frame
- 4. Update tables
- 5. Set validation bit = **v**
- 6. Restart the instruction that caused the page fault



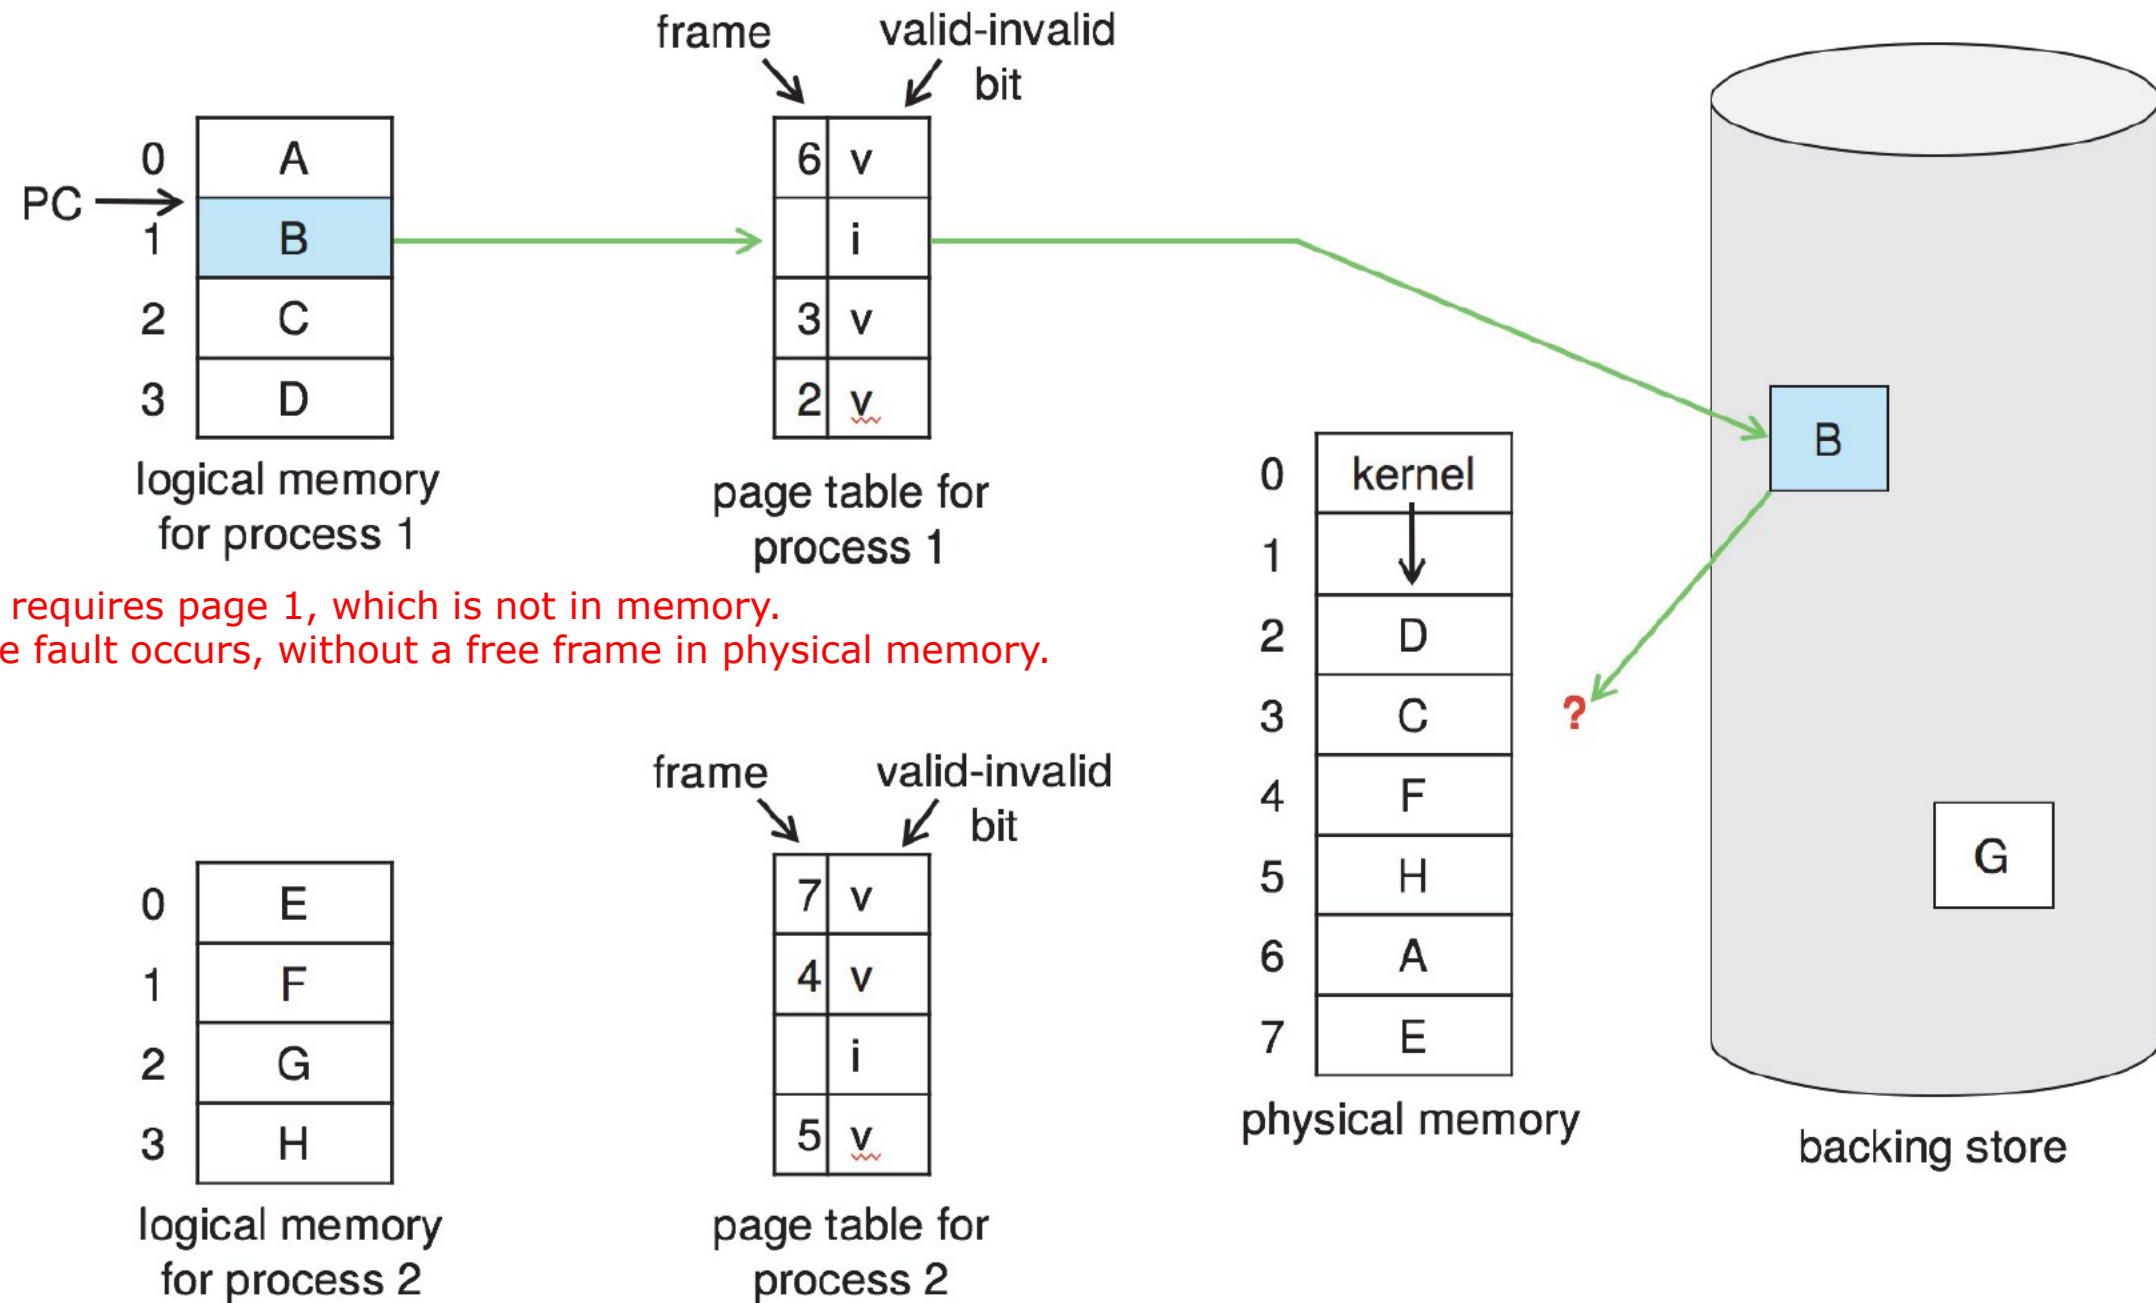


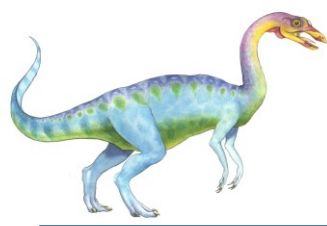
Steps in Handling a Page Fault





Need For Page Replacement

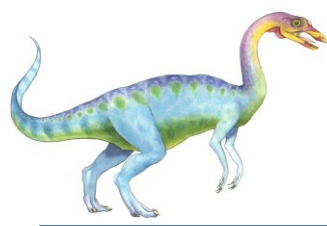




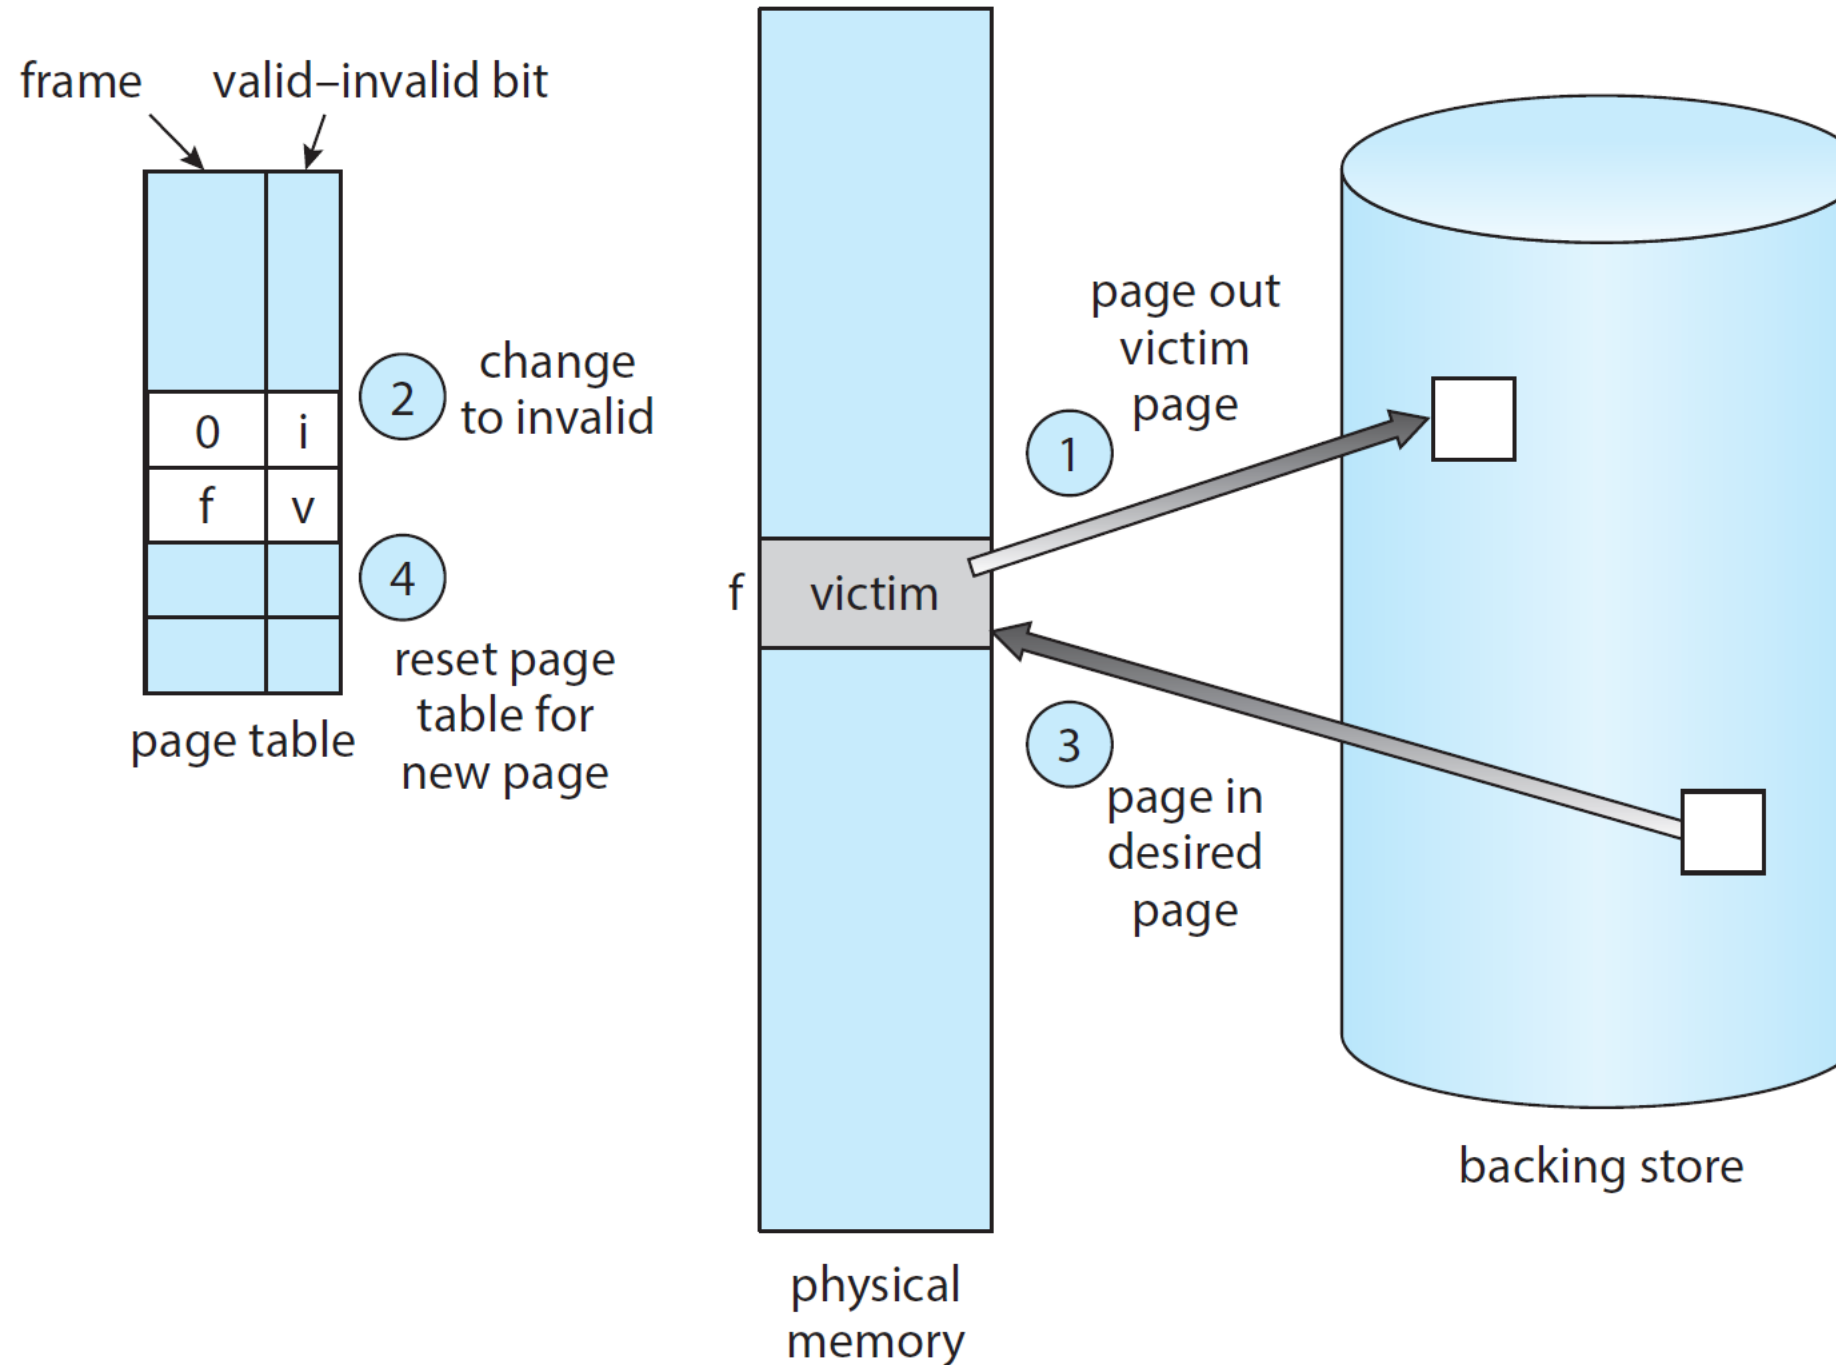
Basic Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**.
 - Write victim frame to disk if dirty.
3. Bring the desired page into the (newly) free frame; update the page and frame tables.
4. Continue the process by restarting the instruction that caused the trap.





Page Replacement



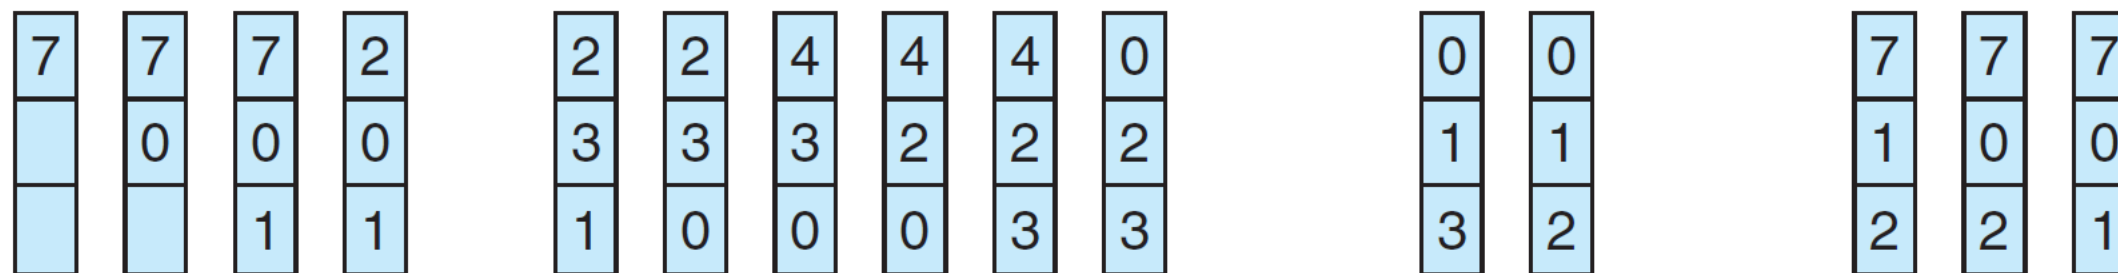


First-In-First-Out (FIFO) Algorithm

- Use a FIFO queue. When a page is brought into memory, append it into the queue.
- When a page must be replaced, choose the oldest page in memory (the first one in the queue).
- Example: Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)

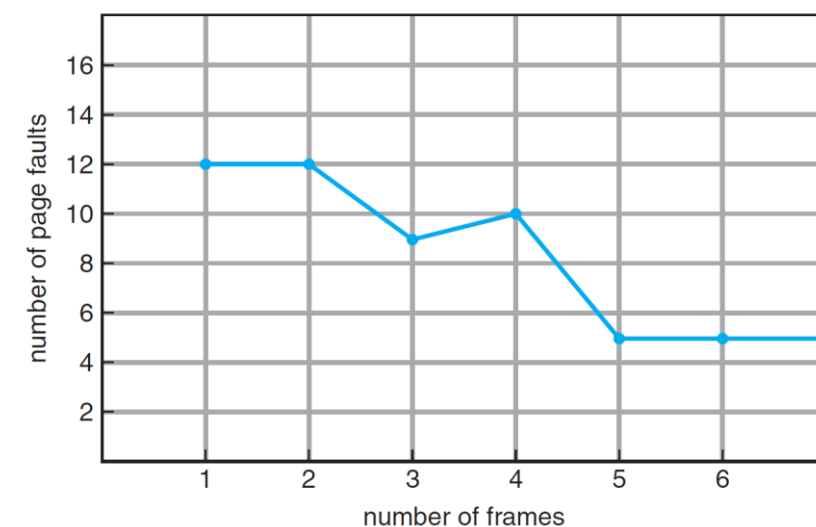
reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

- 15 page faults
- Adding more frames can cause more page faults!
 - Belady's Anomaly**
 - consider reference string: 1,2,3,4,1,2,5,1,2,3,4,5



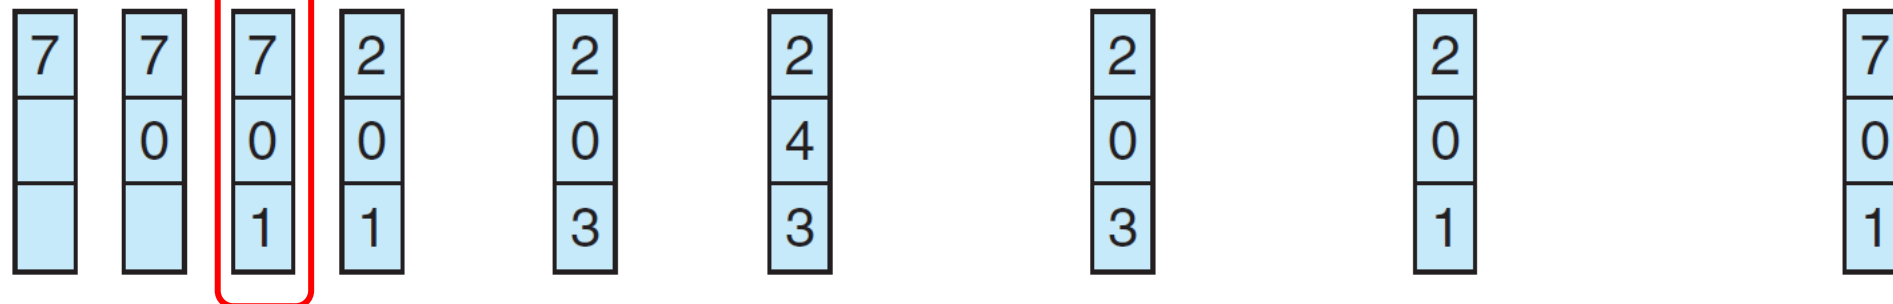


Optimal Algorithm

- ❑ Replace page that will not be used for longest period of time
- ❑ How do you know this?
 - ❑ Can't read the future
- ❑ Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Reference to page 2 causes a page replacement
7 will be used again in 14 references
0 will be used again in 1 reference
1 will be used again in 10 reference
So, replace 7

- ❑ 9 page faults



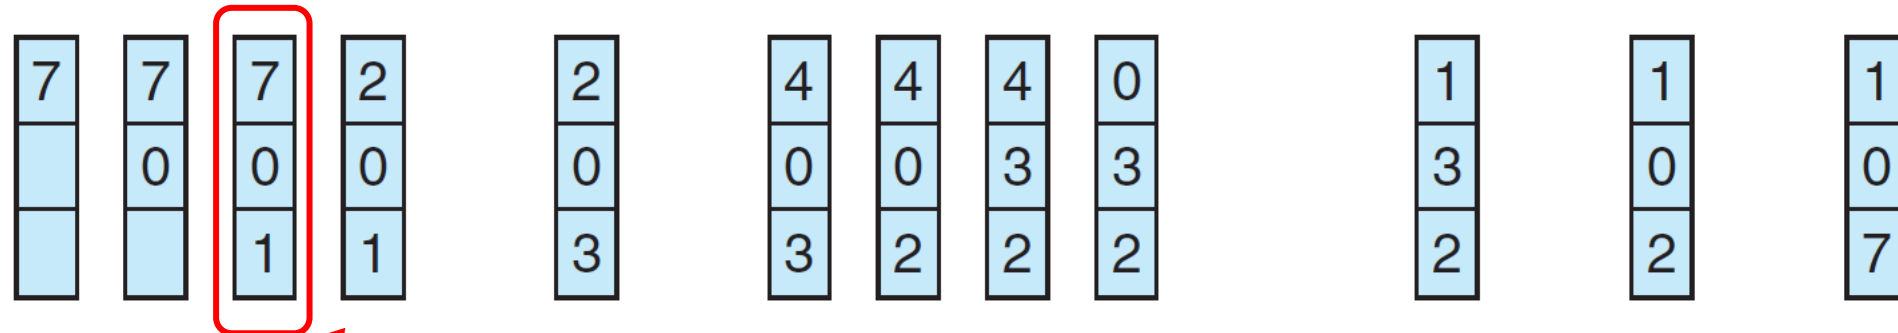


Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Reference to page 2 causes a page replacement
7 has not been used for 3 references
0 has not been used for 2 references
1 has not been used for 1 references
So, replace 7

- 12 faults – better than FIFO but worse than Optimal Algorithm
- Generally good algorithm and frequently used
- But how to implement?

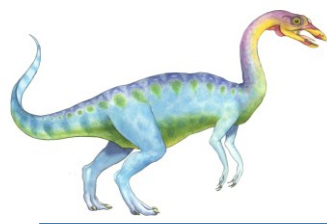




LRU Algorithm (Cont.)

- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to find smallest value
 - ▶ Search through table needed
- Stack implementation
 - Keep a stack of page numbers in a double link form:
 - Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
 - But each update more expensive
 - No search for replacement
- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly

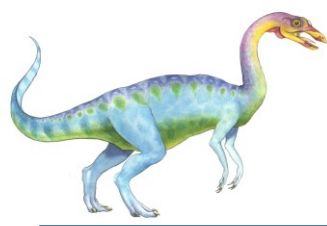




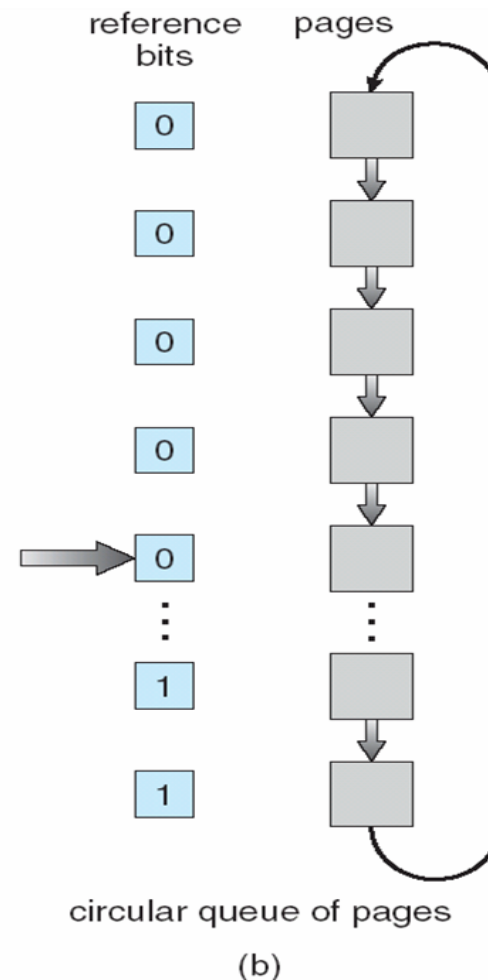
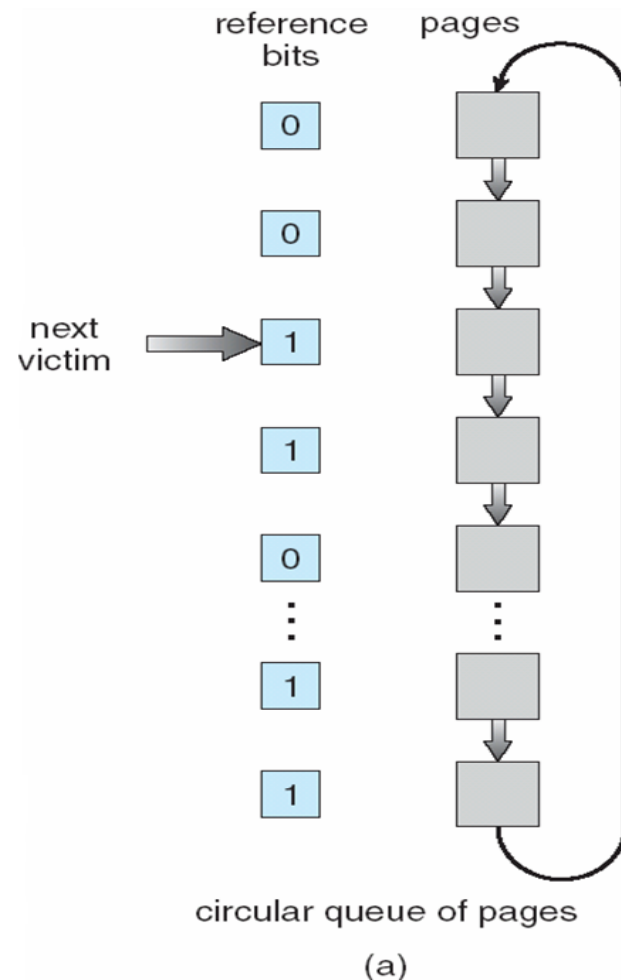
LRU Approximation Algorithms

- LRU needs special hardware and still slow
- **Reference bit**
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace any with reference bit = 0 (if one exists)
 - ▶ We do not know the order, however
- **Second-chance algorithm**
 - Generally FIFO, plus hardware-provided reference bit
 - Clock replacement
 - If page to be replaced has
 - ▶ Reference bit = 0 -> replace it
 - ▶ reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules



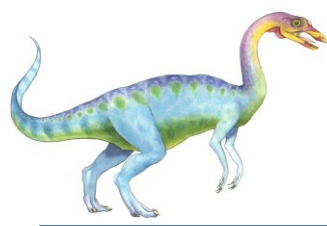


Second-Chance (clock) Page-Replacement Algorithm



- Use a circular queue to keep pages in FIFO manner.
 - Oldest page in memory is at the head of queue and will be replaced first.
 - But if it has been used, it will be given a second chance.
 - Each page has a reference bit. Initialized to zero and set to one when accessed. Also reset to zero periodically.
- If page to be replaced has
 - Reference bit = 0 → replace it
 - reference bit = 1 then:
 - ▶ set reference bit 0, leave page in memory
 - ▶ replace next page, subject to same rules





Shared Pages

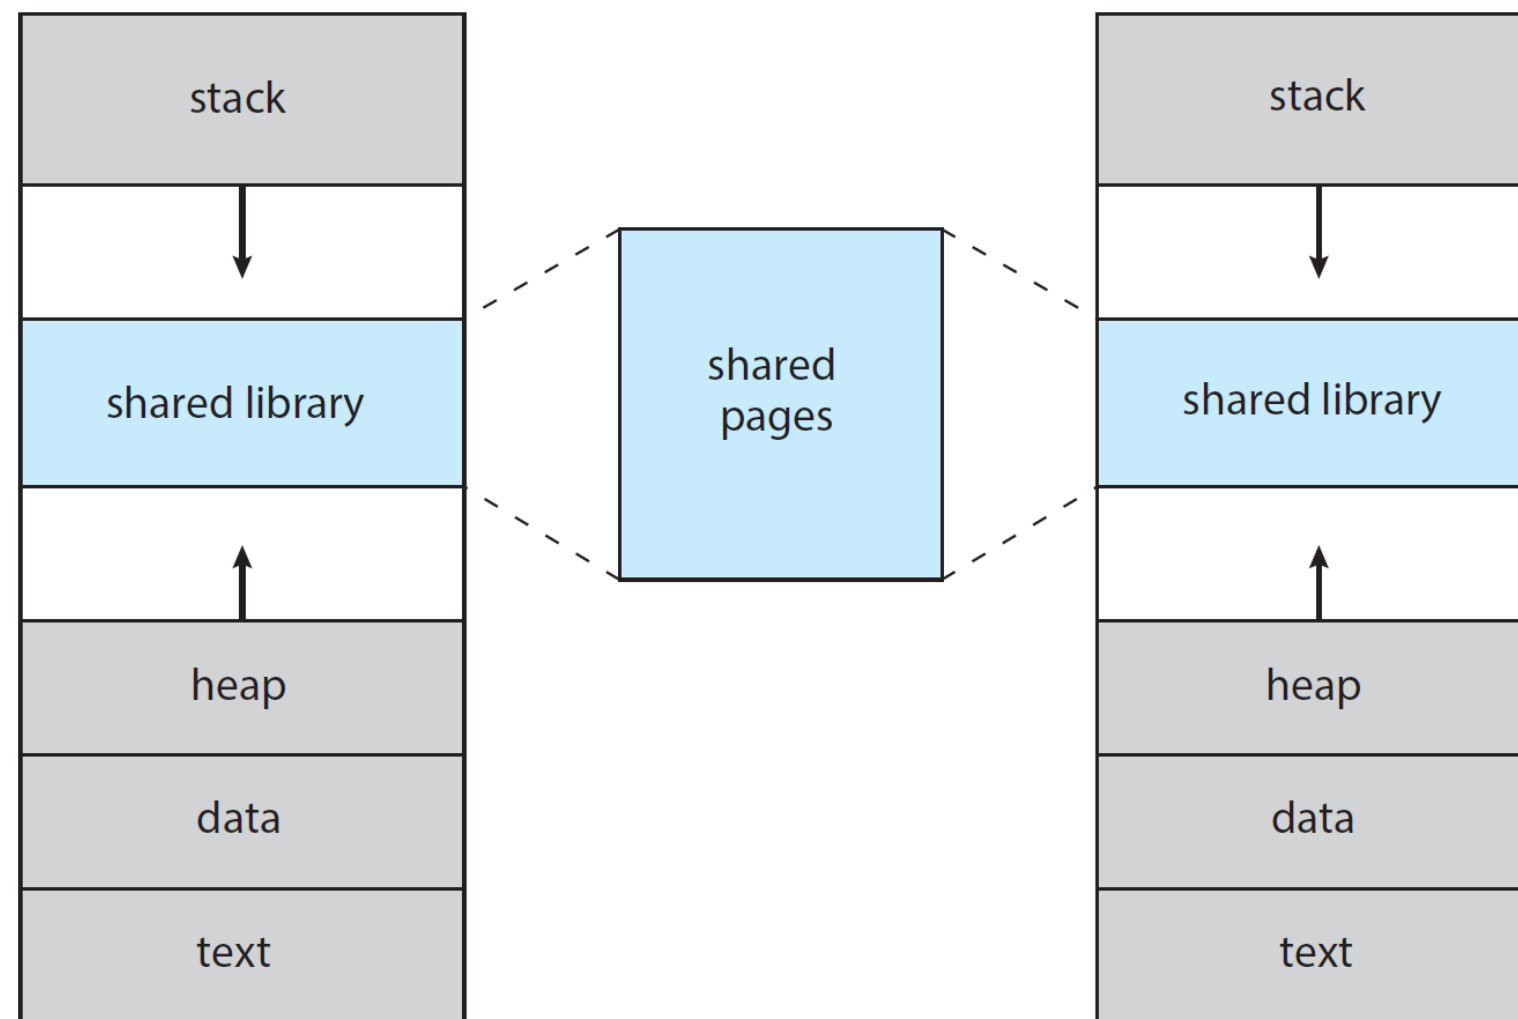
- Virtual memory allows pages to be shared among processes
- Logical memory pages of multiple processes are mapped to the same physical memory frames.
- This technique enables many capabilities:
 - Shared libraries
 - Memory-Mapped files
 - Shared-memory communication





Shared Library Using Virtual Memory

- Code libraries can be shared by several processes.
- Each process considers the library as part of its virtual address space.
- The actual pages where the libraries reside are shared.





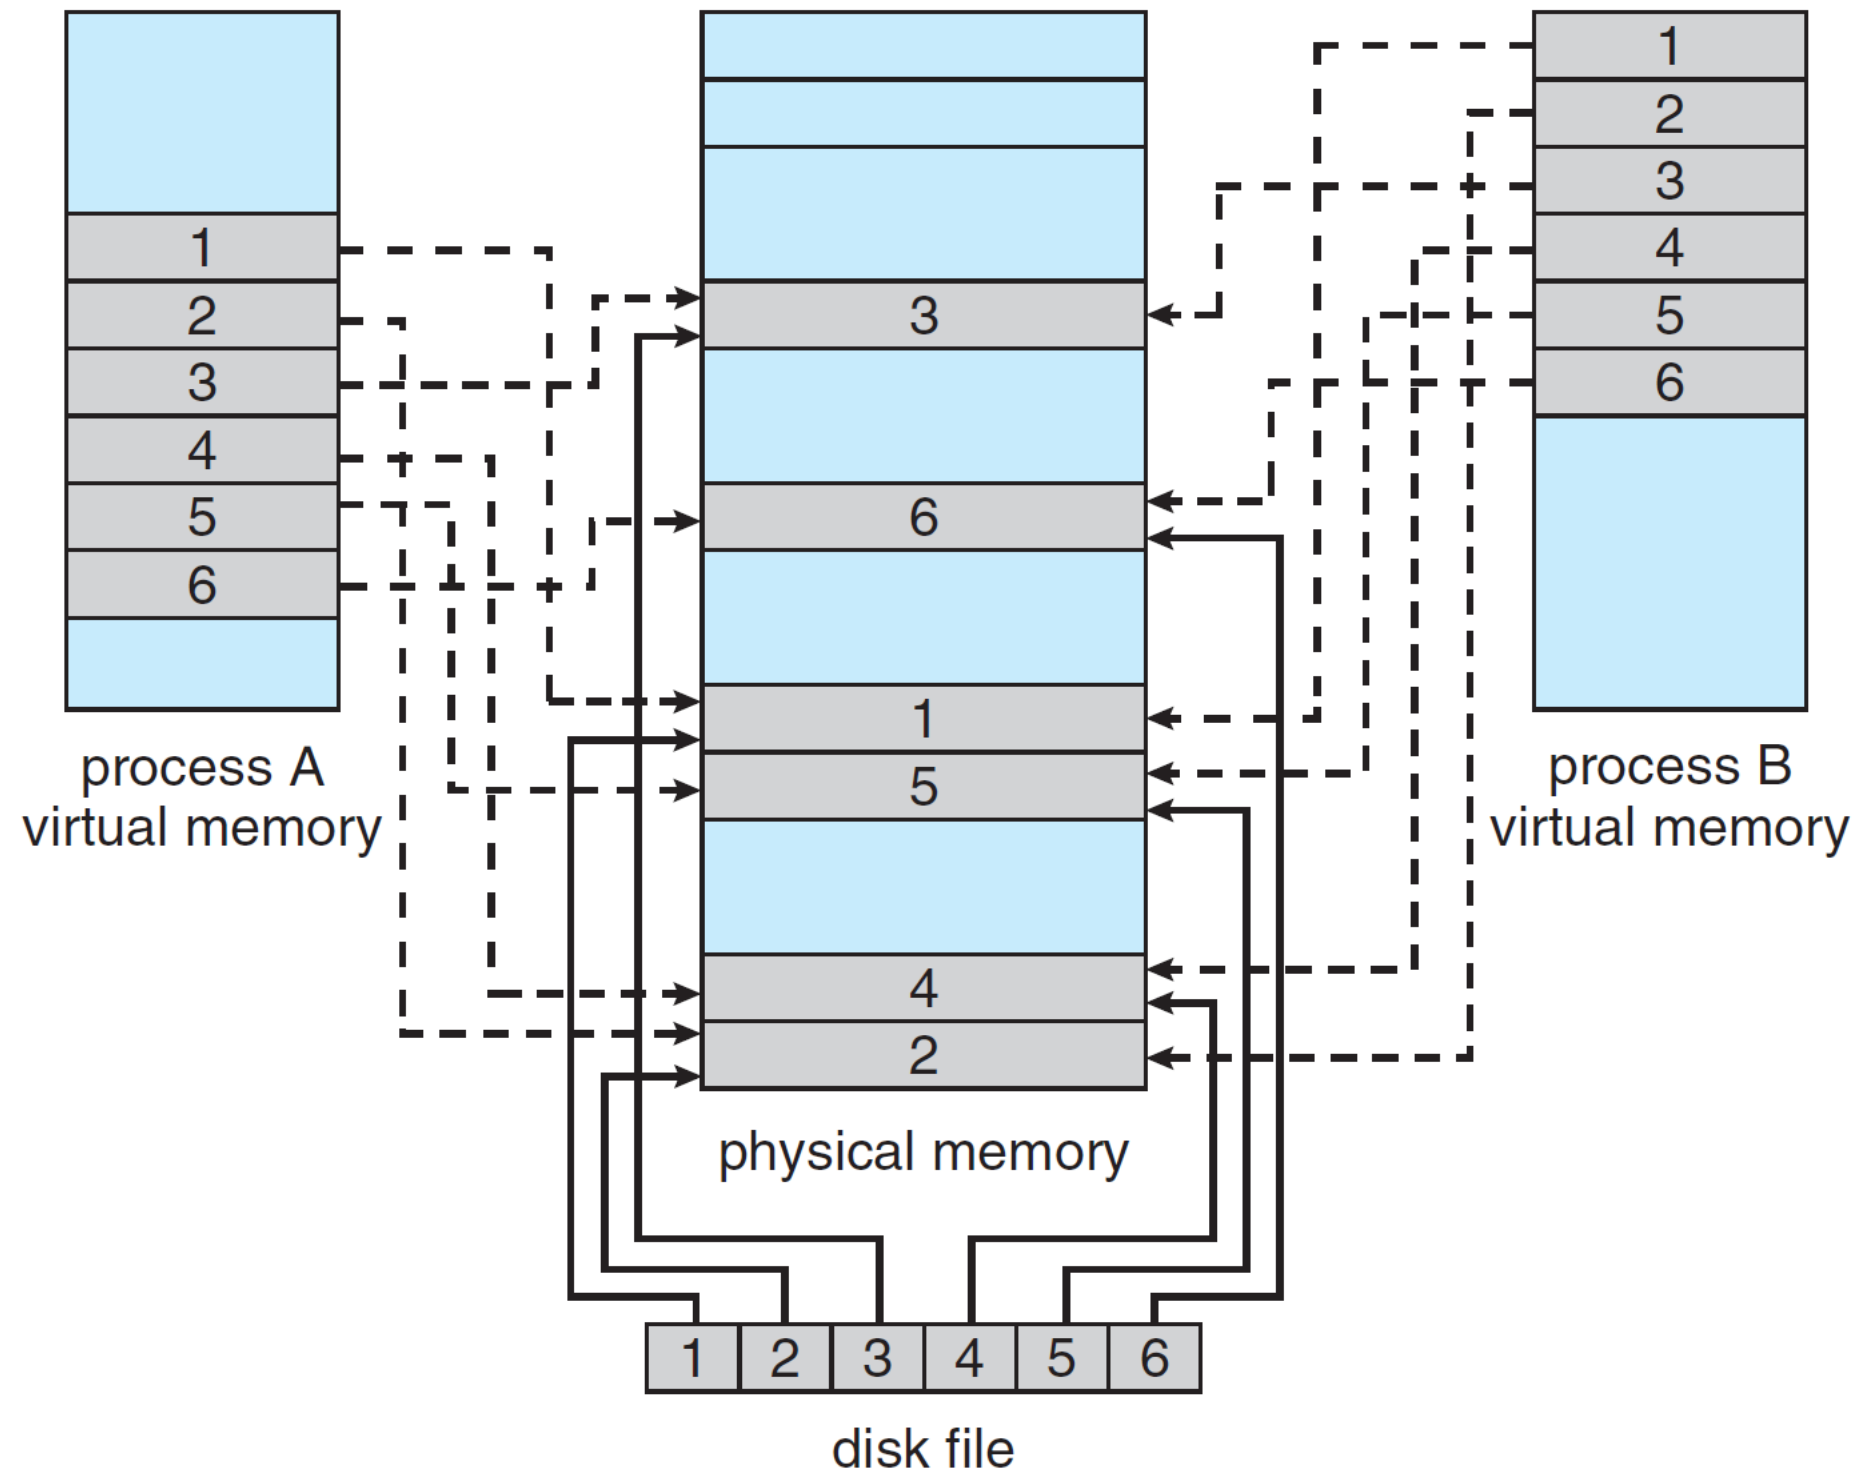
Memory-Mapped Files

- ❑ Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- ❑ A file is initially read using demand paging
 - ❑ A page-sized portion of the file is read from the file system into a physical page
 - ❑ Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- ❑ Simplifies and speeds file access by driving file I/O through memory rather than `read()` and `write()` system calls
- ❑ Also allows several processes to map the same file allowing the pages in memory to be shared
- ❑ But when does written data make it to disk?
 - ❑ Periodically and / or at file `close()` time





Memory Mapped Files





Memory-Mapped Shared Memory

- Processes can communicate through shared memory.
- A process can create a region of memory that it share with another process.
- Processes sharing this region consider it part of their virtual address space, yet the actual physical pages of memory are shared.

